

## Gold Price Forecasting

- The trend of gold price is irregular over all years. This gold price dataset is taken from world gold council and measured in national currency unit per troy ounce. 1 troy ounce is equal to the weight of the gold = 31.1 grams. This dataset contains all the currencies of the world. But, this analysis based on Indian currency.
- ARIMA models are designed to capture the autoregressive (AR) and moving average (MA) components of a time series. The autoregressive component accounts for the relationship between an observation and a certain number of lagged observations, while the moving average component considers the dependency between an observation and a residual error from a regression of previous observations. The "integrated" part of ARIMA represents the differencing operation used to make the time series stationary, meaning that its statistical properties do not vary with time.
- SARIMAX extends ARIMA by incorporating seasonal components. It includes additional parameters to model the seasonal patterns and trends in the data, enabling the capture of repeating patterns that occur over fixed time intervals, such as daily, monthly, or yearly cycles. Additionally, SARIMAX allows for the inclusion of exogenous variables, which are external factors that can affect the time series being analyzed. These variables can provide additional information and improve the accuracy of the forecasting model.
- Prophet has the ability to produce accurate forecasts with minimal effort. It decomposes a time series into three main components: trend, seasonality, and holidays. It automatically detects and handles various types of seasonal patterns, such as daily, weekly, monthly, and yearly seasonality. It can also handle irregular seasonalities. Sometimes, Prophet forecast better than ARIMA and SARIMAX.
- Let's dive into this gold price forecasting and check that which time series forecaster is forecasting best.

## Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

## Loading Data

```
In [2]: data = pd.read_csv("/kaggle/input/gold-price-dataset/Gold Price.csv")
data.head()
```

```
Out[2]:
```

	Name	US dollar	Euro	Japanese yen	Pound sterling	Canadian dollar	Swiss franc	Indian rupee	Chinese renmimbi	US dollar.1	...	Vietnamese dong	Egyptian pound	Korean won	Euro.1	Russian ruble	US dollar.2	South African rand	Chi renmin
0	12/29/1978	78.53	55.93	NaN	64.00	NaN	NaN	NaN	NaN	78.53	...	NaN	NaN	NaN	55.93	NaN	78.53	NaN	1
1	1/1/1979	78.53	55.93	NaN	64.00	NaN	NaN	NaN	NaN	78.53	...	NaN	NaN	NaN	55.93	NaN	78.53	NaN	1
2	1/2/1979	78.80	56.02	132.96	64.43	59.66	90.97	14.66	NaN	78.80	...	NaN	NaN	31.05	56.02	NaN	78.80	11.49	1
3	1/3/1979	75.96	54.68	134.67	62.42	59.81	92.56	14.73	NaN	75.96	...	NaN	NaN	31.21	54.68	NaN	75.96	11.35	1
4	1/4/1979	77.54	55.82	134.53	63.97	59.75	92.69	14.81	NaN	77.54	...	NaN	NaN	31.38	55.82	NaN	77.54	11.40	1

5 rows × 25 columns

```
In [3]: data.shape
```

```
Out[3]: (11581, 25)
```

## Null Data

```
In [4]: nulldata = data.isnull().sum()
nulldata[nulldata>0]
```

```
Out[4]:
```

Japanese yen	2
Canadian dollar	2
Swiss franc	2
Indian rupee	2
Chinese renmimbi	1571
Turkish lira	1045
Saudi riyal	2
Indonesian rupiah	2
UAE dirham	2
Thai baht	2
Vietnamese dong	2851
Egyptian pound	2591
Korean won	2
Russian ruble	3661
South African rand	2
Chinese renmimbi.1	1571
Canadian dollar.1	2
Australian dollar	2

dtype: int64

```
In [5]: data = data.rename(columns = {"Name":"Date"})
data.head()
```

Out[5]:

	Date	US dollar	Euro	Japanese yen	Pound sterling	Canadian dollar	Swiss franc	Indian rupee	Chinese renminbi	US dollar.1	...	Vietnamese dong	Egyptian pound	Korean won	Euro.1	Russian ruble	US dollar.2	South African rand	Chi renmin
0	12/29/1978	78.53	55.93	NaN	64.00	NaN	NaN	NaN	NaN	78.53	...	NaN	NaN	NaN	55.93	NaN	78.53	NaN	1
1	1/1/1979	78.53	55.93	NaN	64.00	NaN	NaN	NaN	NaN	78.53	...	NaN	NaN	NaN	55.93	NaN	78.53	NaN	1
2	1/2/1979	78.80	56.02	132.96	64.43	59.66	90.97	14.66	NaN	78.80	...	NaN	NaN	31.05	56.02	NaN	78.80	11.49	1
3	1/3/1979	75.96	54.68	134.67	62.42	59.81	92.56	14.73	NaN	75.96	...	NaN	NaN	31.21	54.68	NaN	75.96	11.35	1
4	1/4/1979	77.54	55.82	134.53	63.97	59.75	92.69	14.81	NaN	77.54	...	NaN	NaN	31.38	55.82	NaN	77.54	11.40	1

5 rows × 25 columns

## Datetime Extraction

```
In [6]: data['Date'] = pd.to_datetime(data['Date'],format = "%m/%d/%Y")
data['Year'] = data['Date'].dt.year
data['Month'] = data['Date'].dt.month
data['Week'] = data['Date'].dt.week
data['Day'] = data['Date'].dt.day
data.tail()
```

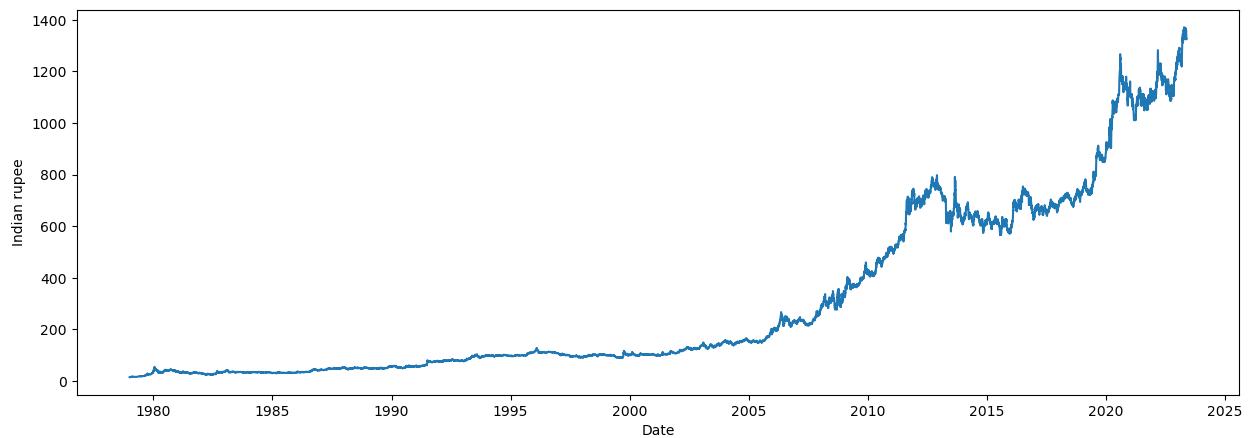
Out[6]:

	Date	US dollar	Euro	Japanese yen	Pound sterling	Canadian dollar	Swiss franc	Indian rupee	Chinese renminbi	US dollar.1	...	Russian ruble	US dollar.2	South African rand	Chinese renminbi.1	Canadian dollar.1	Australian dollar	Year	Mo
11576	2023-05-15	701.84	758.52	847.09	933.43	616.47	457.62	1359.32	589.37	701.84	...	2591.61	701.84	2272.44	589.37	616.47	643.75	2023	
11577	2023-05-16	697.52	753.78	844.10	928.72	611.10	455.10	1349.43	586.75	697.52	...	2610.43	697.52	2261.25	586.75	611.10	642.19	2023	
11578	2023-05-17	686.03	744.52	835.80	915.59	602.36	449.11	1330.05	579.92	686.03	...	2562.67	686.03	2250.99	579.92	602.36	633.61	2023	
11579	2023-05-18	681.13	742.70	837.01	913.37	598.01	449.12	1323.86	578.30	681.13	...	2532.52	681.13	2247.36	578.30	598.01	631.84	2023	
11580	2023-05-19	681.58	741.99	836.91	911.77	599.45	447.68	1325.86	578.21	681.58	...	2534.20	681.58	2256.24	578.21	599.45	628.04	2023	

5 rows × 29 columns

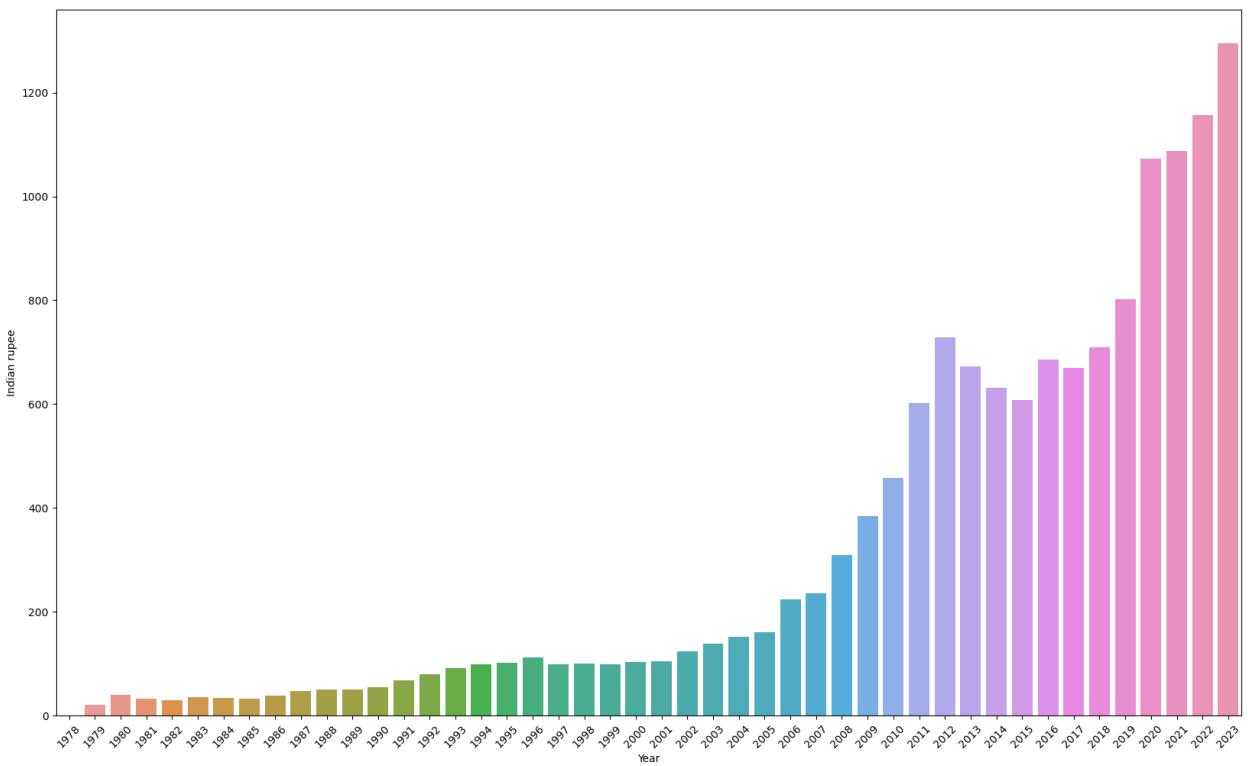
## The Trend

```
In [7]: plt.subplots(figsize = (15,5))
sns.lineplot(data,x = data['Date'], y = data['Indian rupee'])
plt.show()
```



## Visualizing over years

```
In [8]: plt.subplots(figsize = (20,12))
sns.barplot(data,x = data['Year'], y = data['Indian rupee'], ci=None)
plt.xticks(rotation = 45)
plt.show()
```



## Separating the data of 10 years from now

```
In [9]: data10 = data[(data['Year']>2013)]
data10.shape
```

```
Out[9]: (2448, 29)
```

```
In [10]: dataIndia = data10[['Date','Year','Month','Week','Day','Indian rupee']]
dataIndia.shape
```

```
Out[10]: (2448, 6)
```

```
In [11]: dataIndia.head()
```

```
Out[11]:
```

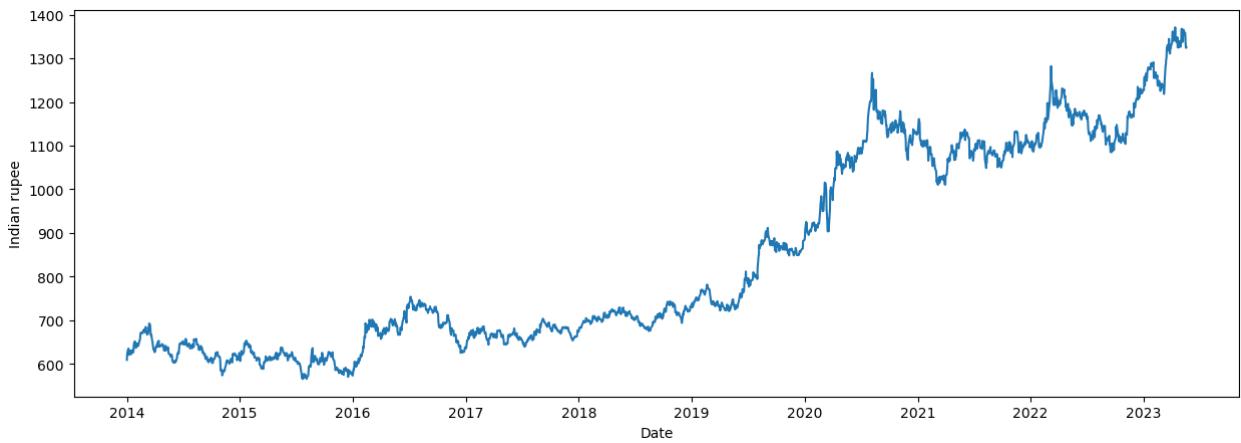
	Date	Year	Month	Week	Day	Indian rupee	
9133	2014-01-01	2014		1	1	1	609.17
9134	2014-01-02	2014		1	1	2	623.05
9135	2014-01-03	2014		1	1	3	627.90
9136	2014-01-06	2014		1	2	6	635.28
9137	2014-01-07	2014		1	2	7	625.47

```
In [12]: dataIndia.isnull().sum()
```

```
Out[12]:
```

Date	0
Year	0
Month	0
Week	0
Day	0
Indian rupee	0
dtype: int64	

```
In [13]: plt.subplots(figsize = (15,5))
sns.lineplot(dataIndia,x = dataIndia['Date'], y = dataIndia['Indian rupee'])
plt.show()
```



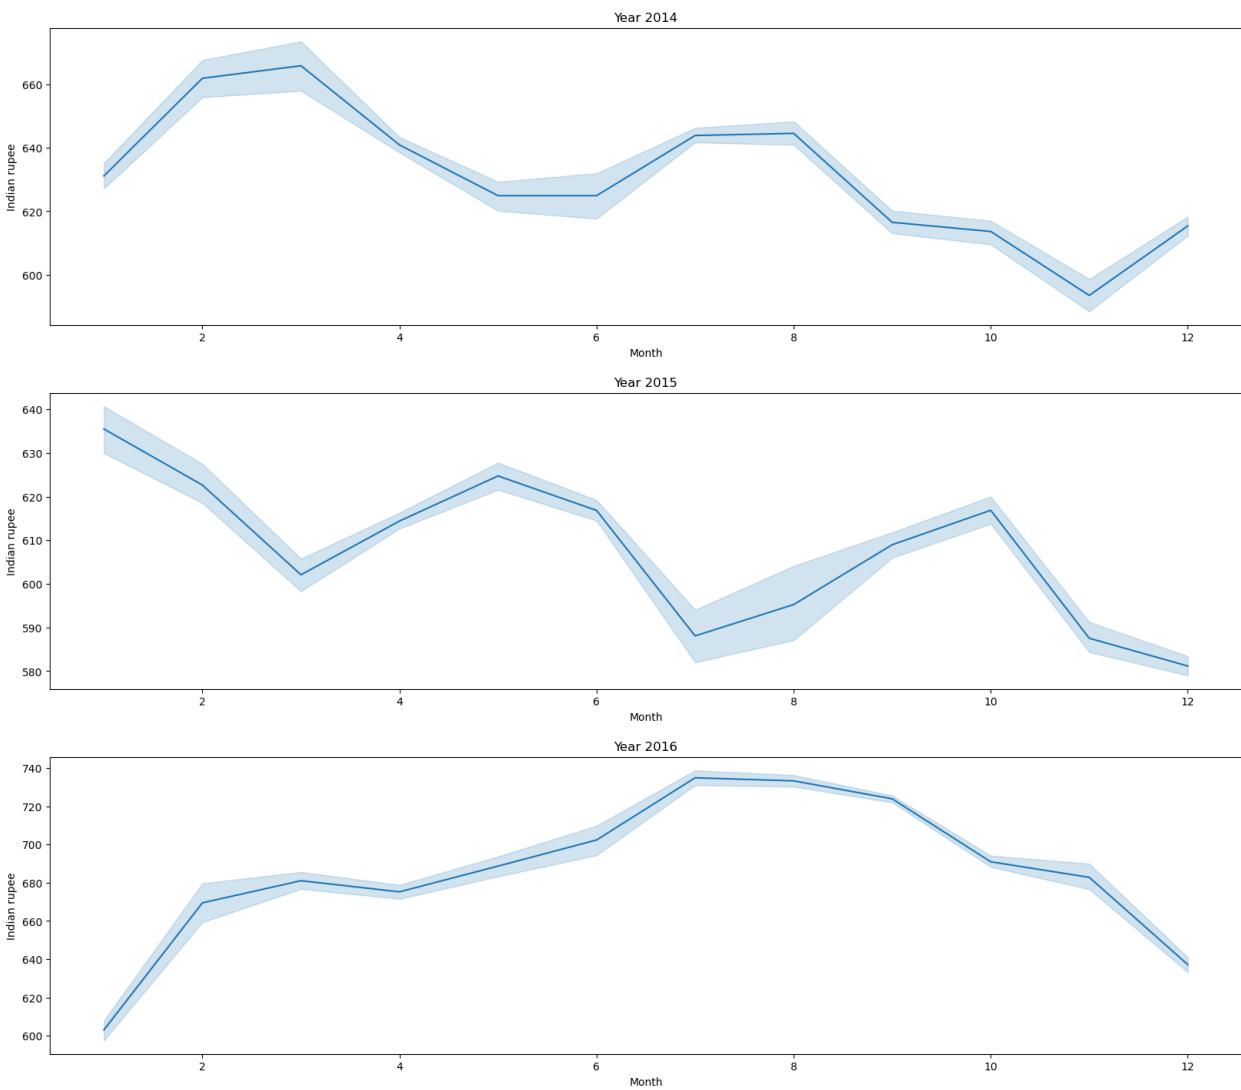
```
In [14]: years = list(dataIndia['Year'].value_counts().index.sort_values())
years
```

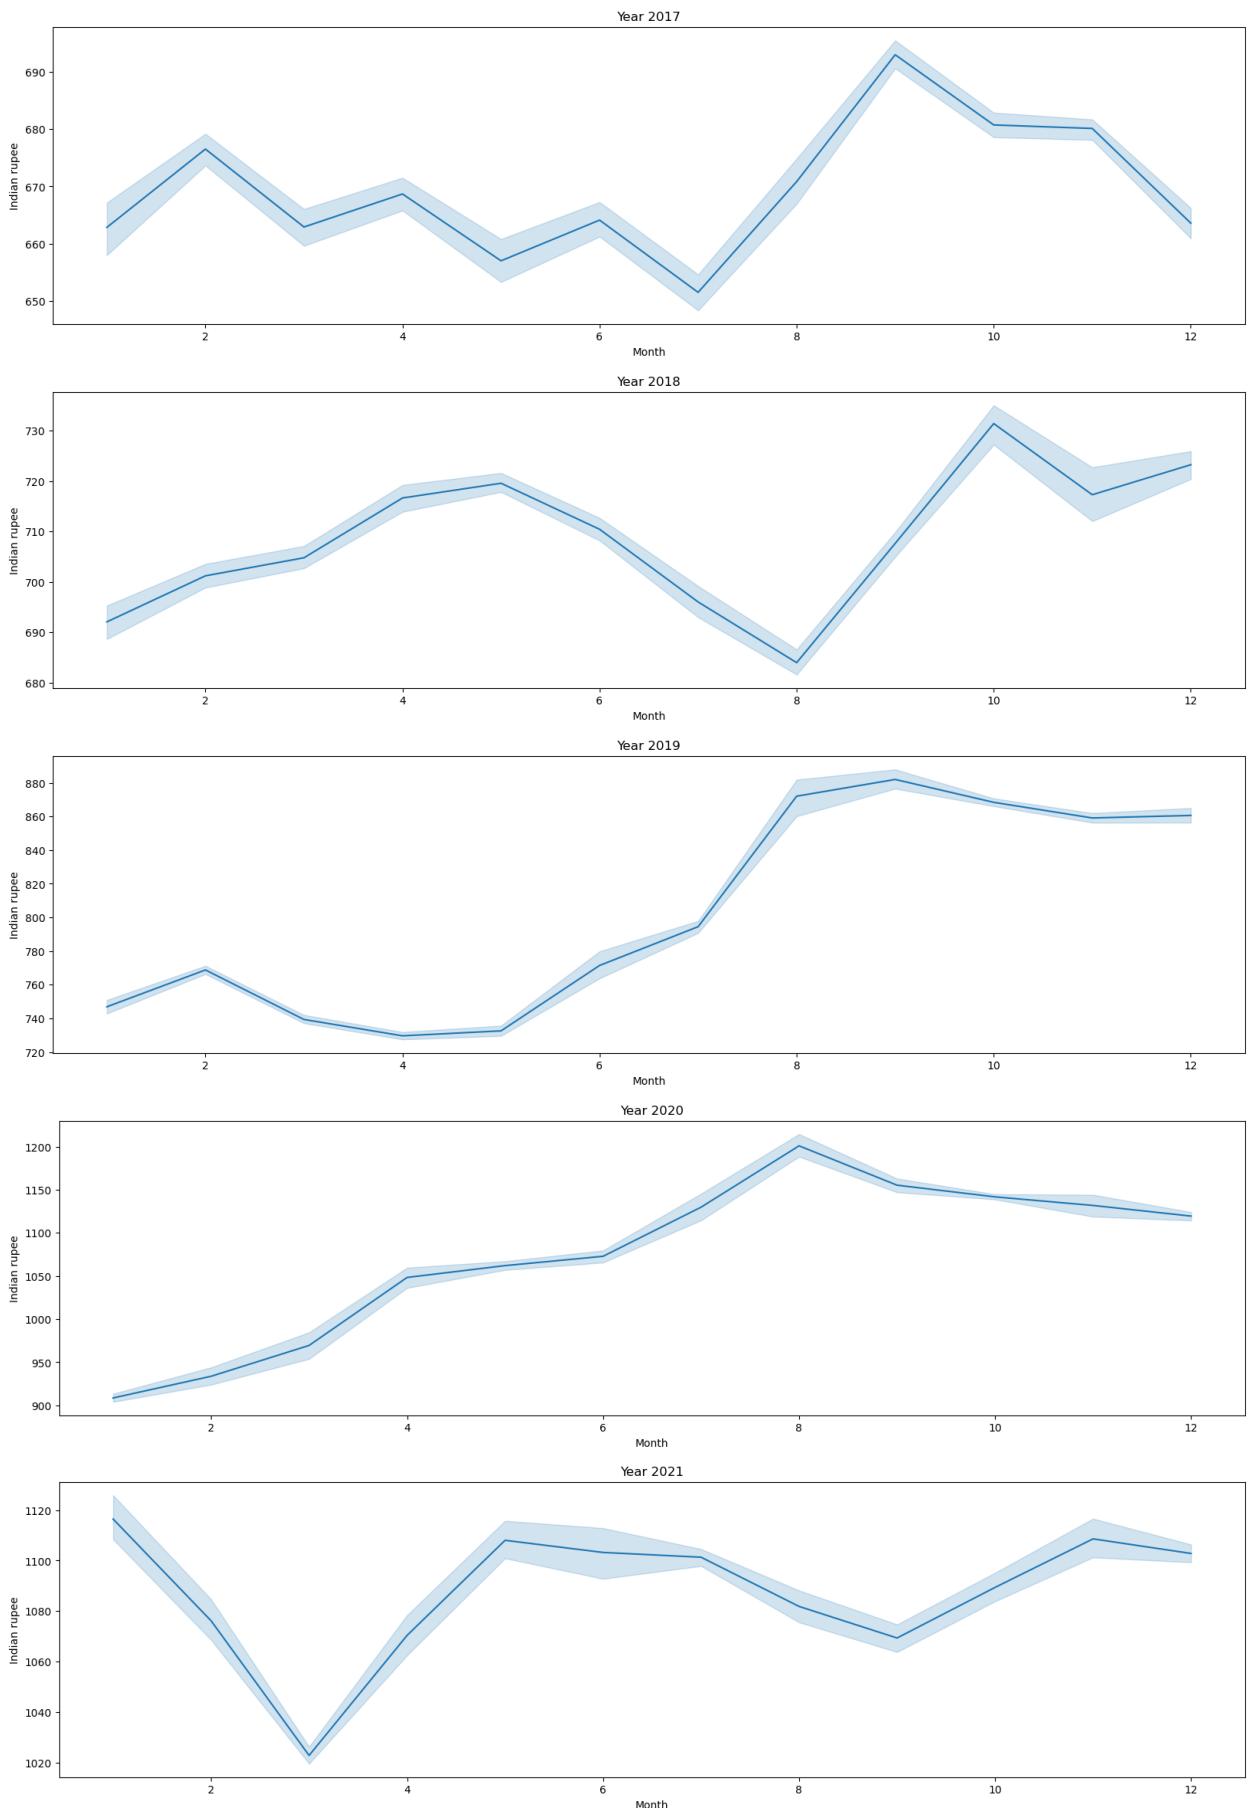
```
Out[14]: [2014, 2015, 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023]
```

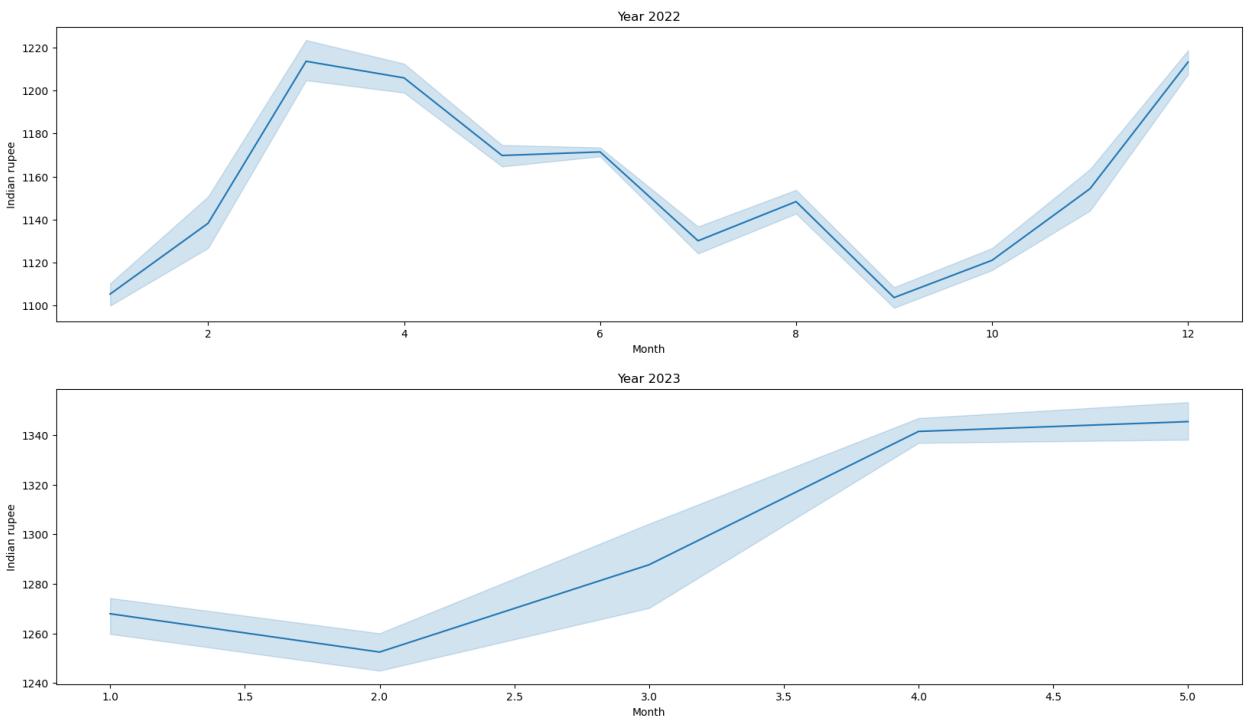
### Helper function for collecting and plotting data of each year after separation

```
In [15]: def Yearly_data(year):
    grp = dataIndia[(dataIndia.Year == year)]
    plt.subplots(figsize = (20,5))
    sns.lineplot(grp,x = grp.Month, y = grp['Indian rupee'])
    plt.title(f'Year {year}')
    plt.show()
```

```
In [16]: for i in range(len(years)):
    Yearly_data(years[i])
```







- It's tough to say which month is best to buy, but if we go with the majority, **April to August** will be the best month for buying gold as found from the past 10 years data.

```
In [17]: dataIndia.head()
```

```
Out[17]:
```

	Date	Year	Month	Week	Day	Indian rupee
9133	2014-01-01	2014	1	1	1	609.17
9134	2014-01-02	2014	1	1	2	623.05
9135	2014-01-03	2014	1	1	3	627.90
9136	2014-01-06	2014	1	2	6	635.28
9137	2014-01-07	2014	1	2	7	625.47

## Let's try ARIMA

```
In [18]: dataIndia.set_index(dataIndia['Date'], inplace = True)
dataIndia = dataIndia.drop(columns = 'Date')
dataIndia.tail()
```

```
Out[18]:
```

	Year	Month	Week	Day	Indian rupee
2023-05-15	2023	5	20	15	1359.32
2023-05-16	2023	5	20	16	1349.43
2023-05-17	2023	5	20	17	1330.05
2023-05-18	2023	5	20	18	1323.86
2023-05-19	2023	5	20	19	1325.86

## Helper function for checking stationarity

```
In [19]: def check_stationarity(series):
    from statsmodels.tsa.stattools import adfuller
    result = adfuller(series)
    print('ADF Statistic: %f' % result[0])
    print('p-value: %f' % result[1])
    print('Critical Values:')
    for key, value in result[4].items():
        print(' %s: %.3f' % (key, value))
    if (result[1] <= 0.05) & (result[4]['5%'] > result[0]):
        print("\u001b[32mStationary\u001b[0m")
    else:
        print("\u001b[31mNon-stationary\u001b[0m")
```

```
In [20]: from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf
```

```
In [21]: check_stationarity(dataIndia['Indian rupee'])
```

```
ADF Statistic: 0.327753
```

```
p-value: 0.978583
```

```
Critical Values:
```

```
1%: -3.433
```

```
5%: -2.863
```

```
10%: -2.567
```

```
Non-stationary
```

## Try differencing

```
In [22]: dataIndia['INR_diff'] = dataIndia['Indian rupee'].diff().fillna(0)
```

```
check_stationarity(dataIndia['INR_diff'])
```

```
ADF Statistic: -14.456611
```

```
p-value: 0.000000
```

```
Critical Values:
```

```
1%: -3.433
```

```
5%: -2.863
```

```
10%: -2.567
```

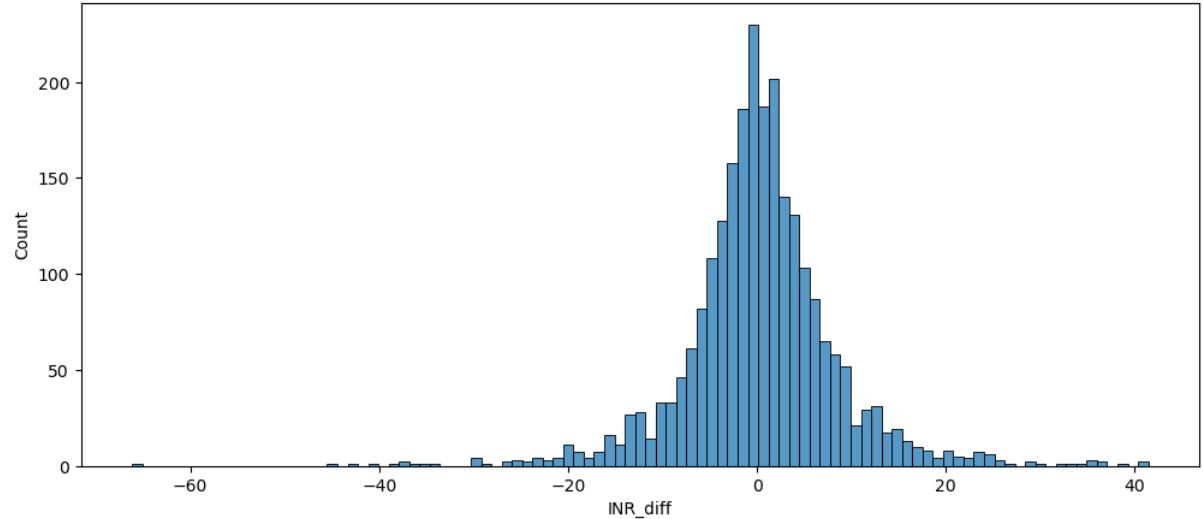
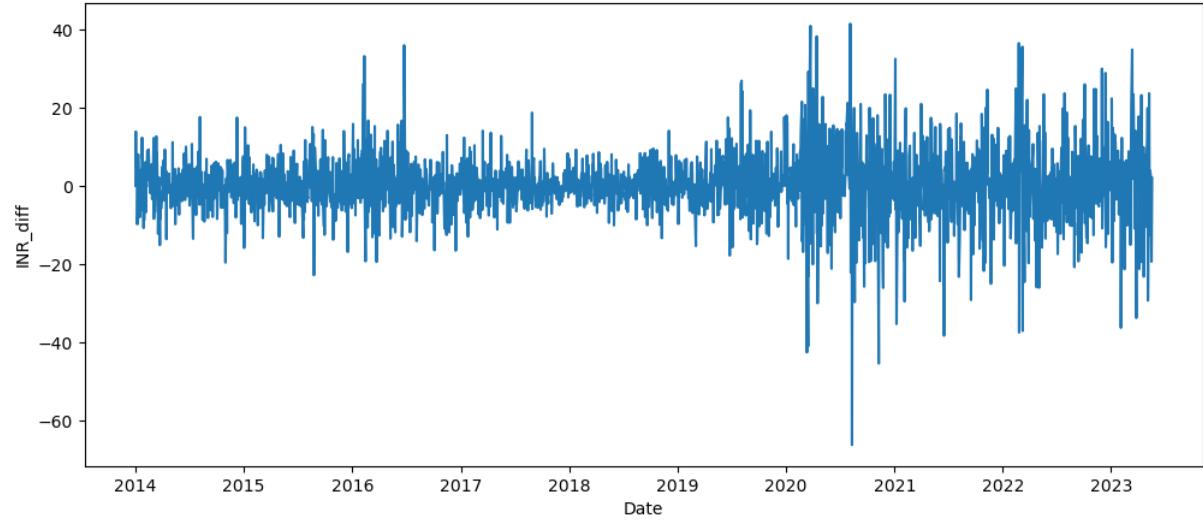
```
Stationary
```

```
In [23]: f, ax = plt.subplots(nrows=1, ncols=1, figsize=(12,5))
```

```
sns.lineplot(data=dataIndia, x=dataIndia.index, y='INR_diff')  
plt.show()
```

```
f, ax = plt.subplots(nrows=1, ncols=1, figsize=(12,5))
```

```
sns.histplot(dataIndia['INR_diff'])  
plt.show()
```



```
In [24]: train = dataIndia[dataIndia.index < pd.to_datetime("2023-01-01",format = "%Y-%m-%d")]
```

```
test = dataIndia[dataIndia.index > pd.to_datetime("2023-01-01",format = "%Y-%m-%d")]
```

```
train = train[['INR_diff']]
```

```
test = test[['INR_diff']]
```

```
sns.set()
```

```
plt.subplots(figsize = (15,10))
```

```
plt.plot(train,color ='black')
```

```
plt.plot(test,color = 'red')
```

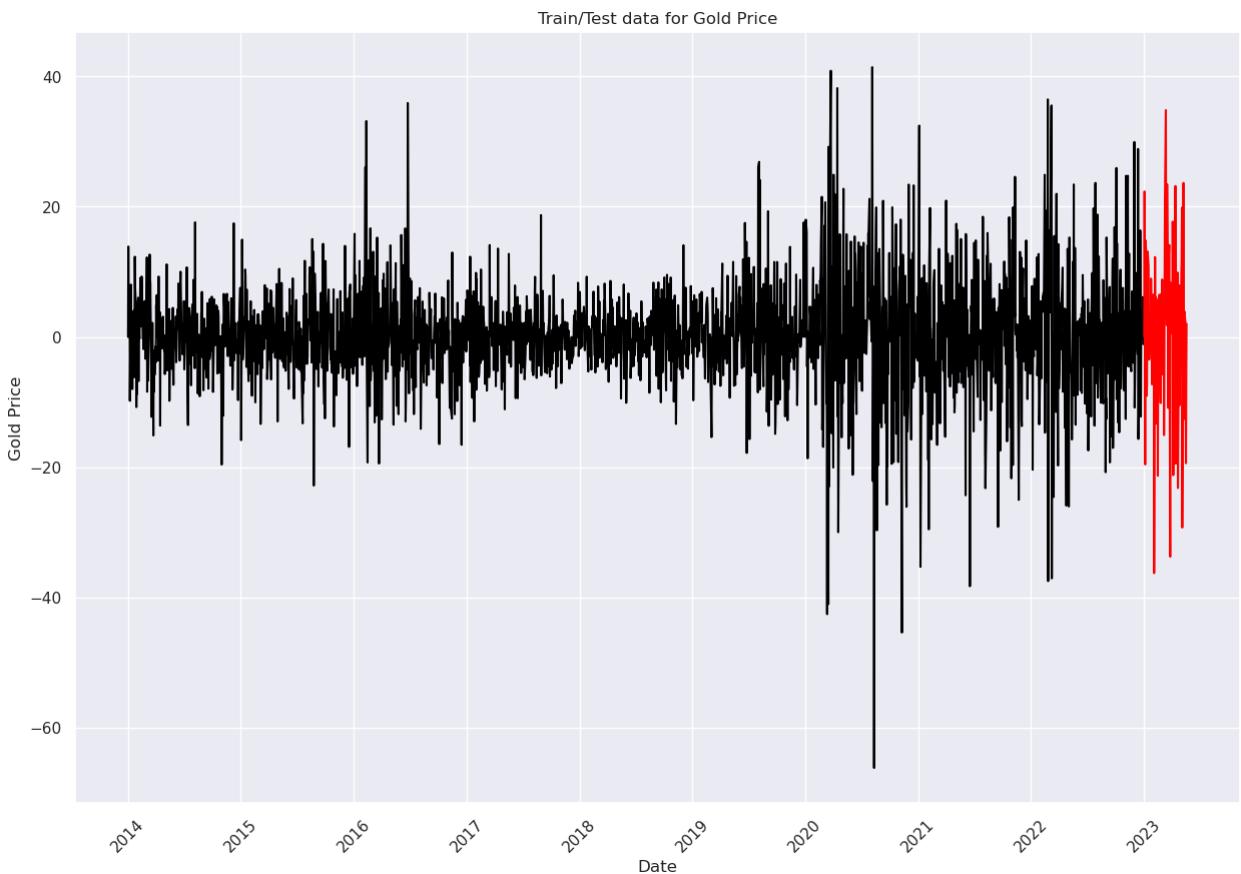
```
plt.ylabel("Gold Price")
```

```
plt.xlabel("Date")
```

```
plt.xticks(rotation = 45)
```

```
plt.title("Train/Test data for Gold Price")
```

```
plt.show()
```



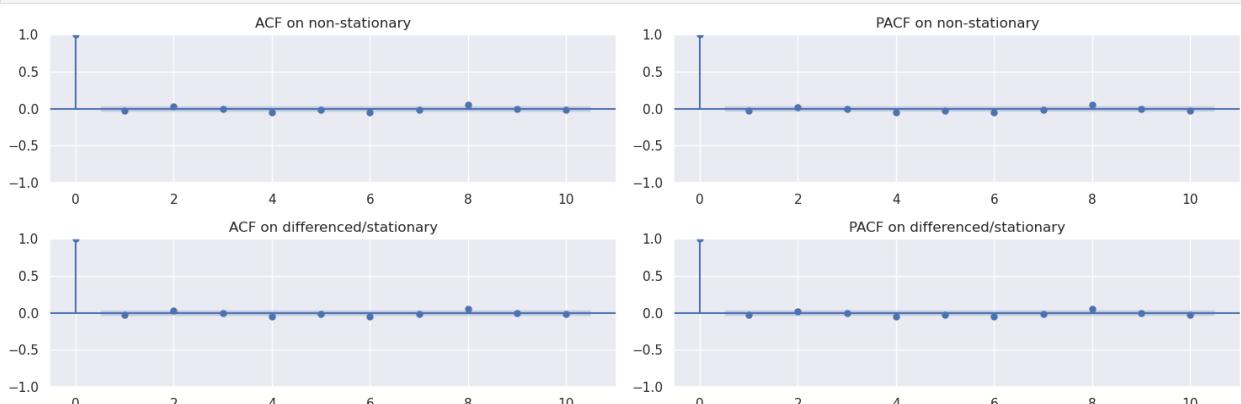
## ACF and PACF

```
In [25]: f, ax = plt.subplots(nrows=2, ncols=2, figsize=(15,5))

plot_acf(train['INR_diff'], lags=10, ax=ax[0, 0], title='ACF on non-stationary')
plot_pacf(train['INR_diff'], lags=10, ax=ax[0, 1], method='ols', title='PACF on non-stationary')

plot_acf(train['INR_diff'], lags=10, ax=ax[1, 0], title='ACF on differenced/stationary')
plot_pacf(train['INR_diff'], lags=10, ax=ax[1, 1], method='ols', title='PACF on differenced/stationary')

plt.tight_layout()
plt.show()
```



- Differencing of 1 will give perfect results.

## Helper function for checking order

```
In [26]: def check_order(p,d,q):
    y = train['INR_diff']
    from statsmodels.tsa.arima.model import ARIMA
    ARIMAm model = ARIMA(y, order = (p,d,q))
    ARIMAm model = ARIMAm model.fit()
    print(ARIMAm model.summary(), "\n")

    y_pred = ARIMAm model.get_forecast(len(test.index))
    y_pred_df = y_pred.conf_int(alpha = 0.05)
    y_pred_df['Predictions'] = ARIMAm model.predict(start = y_pred_df.index[0], end = y_pred_df.index[-1])
    y_pred_df.index = test.index
    y_pred_out = y_pred_df['Predictions']
```

```

from sklearn.metrics import mean_squared_error
arima_rmse = np.sqrt(mean_squared_error(test['INR_diff'].values, y_pred_df["Predictions"]))
print("RMSE: ",arima_rmse, '\n')

ARIMAmode.plot_diagnostics(figsize = (20,12))

sns.set()
plt.subplots(figsize = (15,5))
plt.plot(train,color = 'black',label = 'Train')
plt.plot(test,color = 'red',label = 'Test')
plt.plot(y_pred_out, color='Yellow', label = 'ARIMA Predictions')
plt.ylabel("Gold Price")
plt.xlabel("Date")
plt.xticks(rotation = 45)
plt.title("Train/Test and predictions for Gold Price")
plt.tight_layout()
plt.show()

```

In [27]: check\_order(0,0,0)

```

/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
self._init_dates(dates, freq)
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
self._init_dates(dates, freq)
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
self._init_dates(dates, freq)

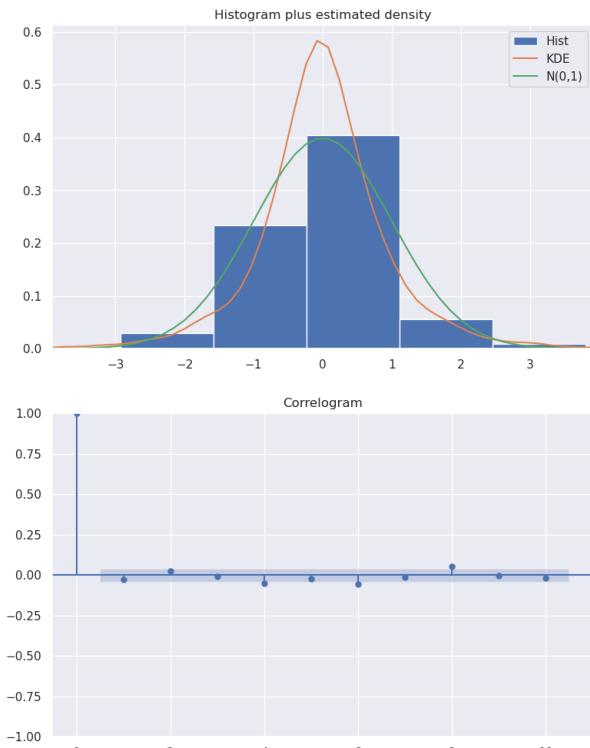
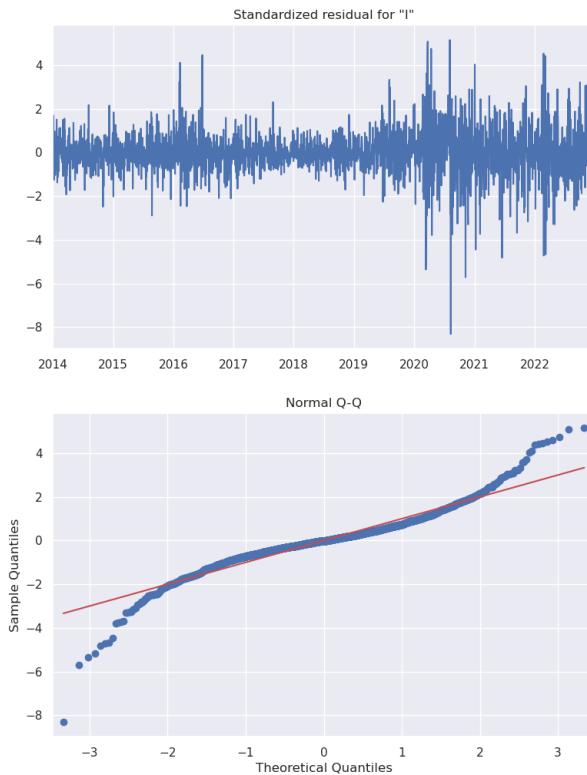
SARIMAX Results
=====
Dep. Variable: INR_diff No. Observations: 2348
Model: ARIMA Log Likelihood: -8211.769
Date: Mon, 29 May 2023 AIC: 16427.537
Time: 14:02:57 BIC: 16439.060
Sample: 01-01-2014 HQIC: 16431.734
- 12-30-2022
Covariance Type: opg
=====

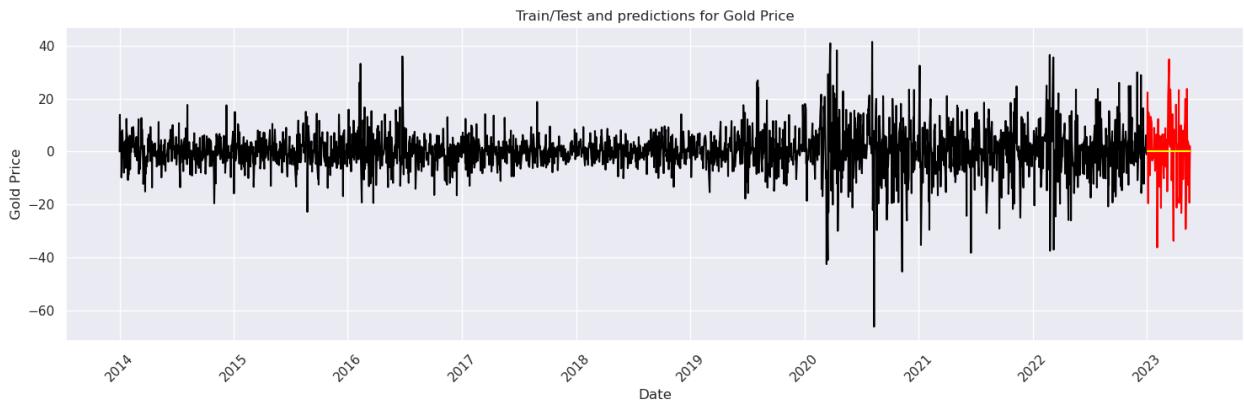
            coef    std err      z   P>|z|      [0.025]     [0.975]
-----+
const    0.2631    0.165    1.591    0.112    -0.061     0.587
sigma2   63.8632   0.924   69.133    0.000    62.053    65.674
-----+
Ljung-Box (L1) (Q): 2.05 Jarque-Bera (JB): 3753.84
Prob(Q): 0.15 Prob(JB): 0.00
Heteroskedasticity (H): 3.55 Skew: -0.20
Prob(H) (two-sided): 0.00 Kurtosis: 9.18
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```

RMSE: 12.136689951443172





In [28]: `check_order(0,1,0)`

```
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)

SARIMAX Results
=====
Dep. Variable:      INR_diff    No. Observations:          2348
Model:             ARIMA(0, 1, 0)   Log Likelihood:       -9056.354
Date:        Mon, 29 May 2023   AIC:                  18114.709
Time:            14:03:01     BIC:                  18120.470
Sample:         01-01-2014   HQIC:                 18116.807
                    - 12-30-2022
Covariance Type:            opg

=====
```

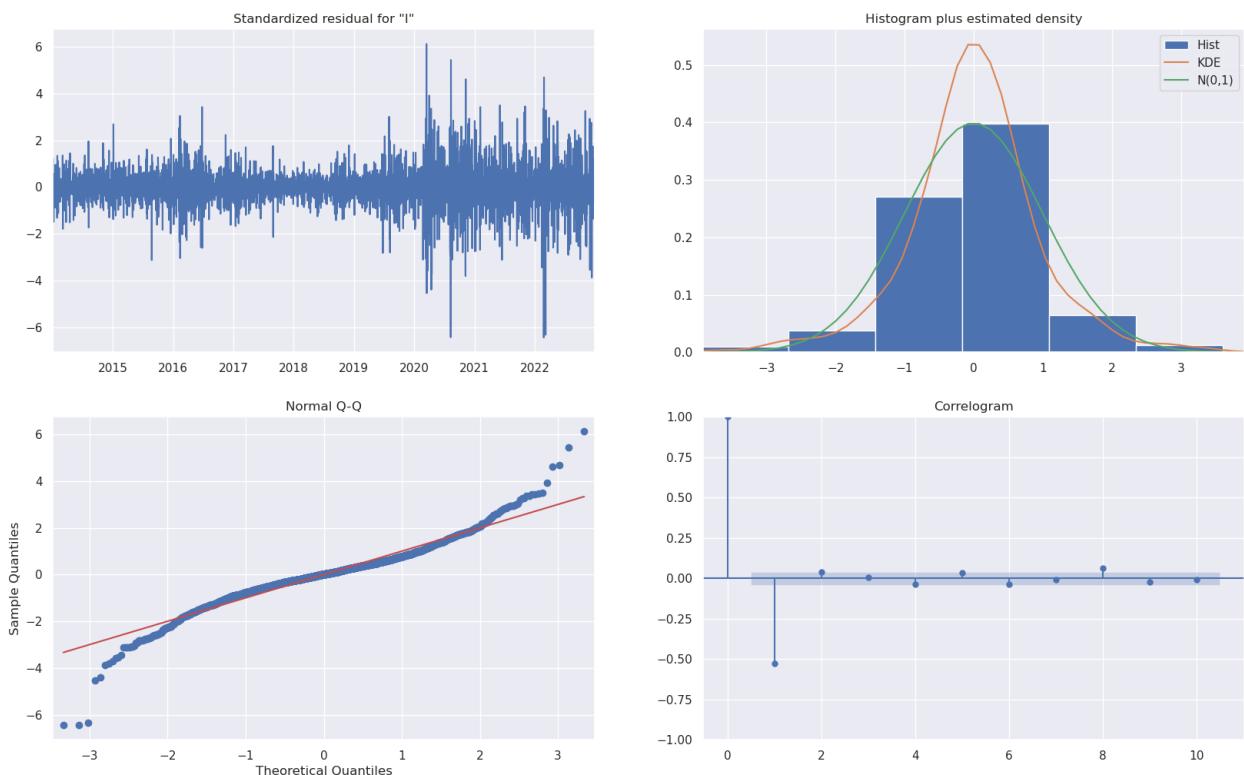
coef	std err	z	P> z	[0.025	0.975]	
sigma2	131.5669	2.096	62.767	0.000	127.459	135.675

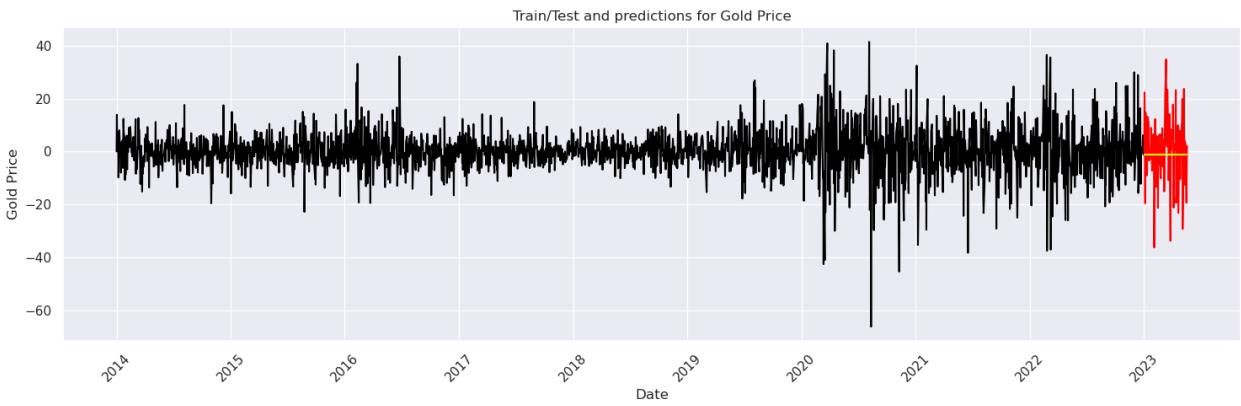
```
=====
Ljung-Box (Q):      650.50   Jarque-Bera (JB):       2190.15
Prob(Q):           0.00   Prob(JB):                   0.00
Heteroskedasticity (H): 3.53   Skew:                  -0.21
Prob(H) (two-sided): 0.00   Kurtosis:                 7.71
=====
```

Warnings:

- [1] Covariance matrix calculated using the outer product of gradients (complex-step).

RMSE: 12.29554964204532





In [29]: `check_order(2,1,2)`

```
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)

SARIMAX Results
=====
Dep. Variable:      INR_diff    No. Observations:          2348
Model:             ARIMA(2, 1, 2)    Log Likelihood       -8211.424
Date:        Mon, 29 May 2023   AIC                  16432.848
Time:                14:03:04   BIC                  16461.652
Sample:         01-01-2014   HQIC                 16443.339
                    - 12-30-2022
Covariance Type:    opg

=====
```

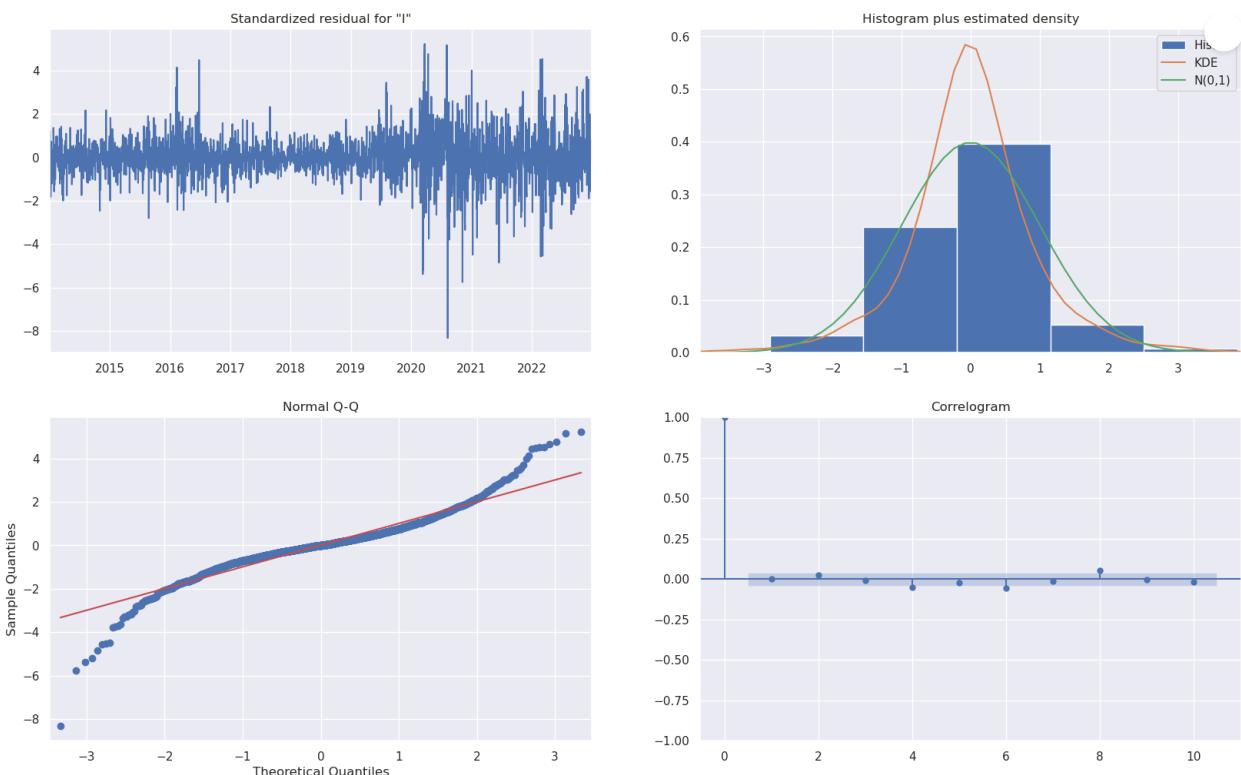
coef	std err	z	P> z	[0.025	0.975]	
ar.L1	-1.0286	0.015	-70.860	0.000	-1.057	-1.000
ar.L2	-0.0295	0.015	-2.038	0.042	-0.058	-0.001
ma.L1	-7.11e-06	375.996	-1.89e-08	1.000	-736.938	736.938
ma.L2	-1.0000	376.016	-0.003	0.998	-737.978	735.978
sigma2	63.8080	2.4e+04	0.003	0.998	-4.7e+04	4.71e+04

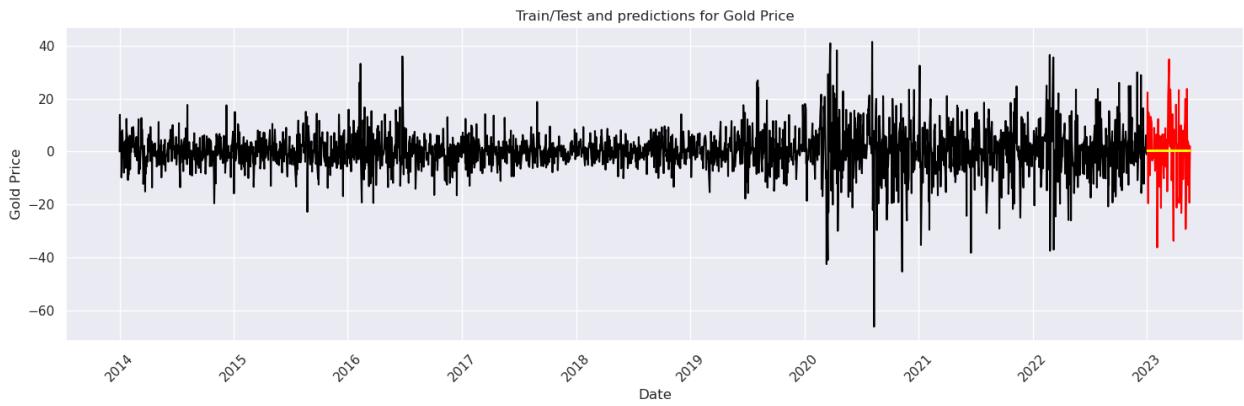
```
Ljung-Box (L1) (Q):      0.00   Jarque-Bera (JB):      3835.65
Prob(Q):                  0.98   Prob(JB):                  0.00
Heteroskedasticity (H):  3.56   Skew:                  -0.21
Prob(H) (two-sided):     0.00   Kurtosis:                 9.25
```

#### Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

RMSE: 12.137488093483809





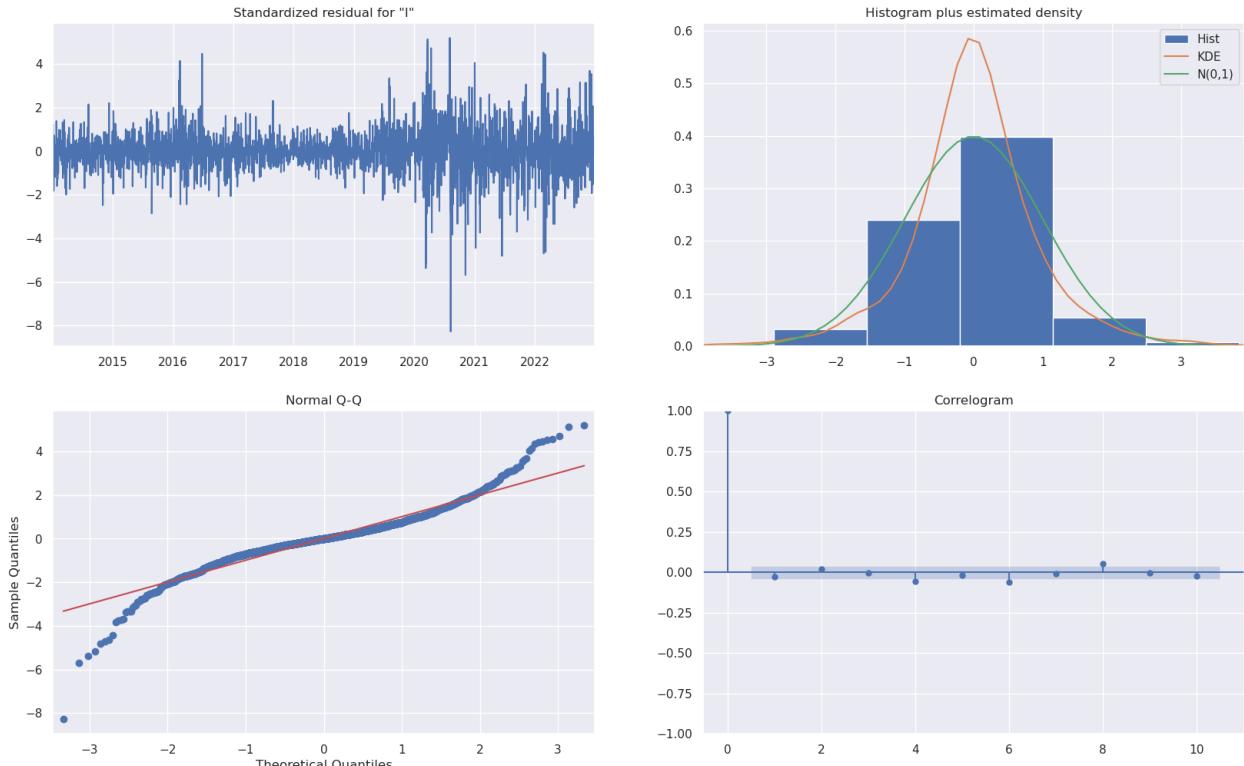
In [30]: `check_order(1,1,2)`

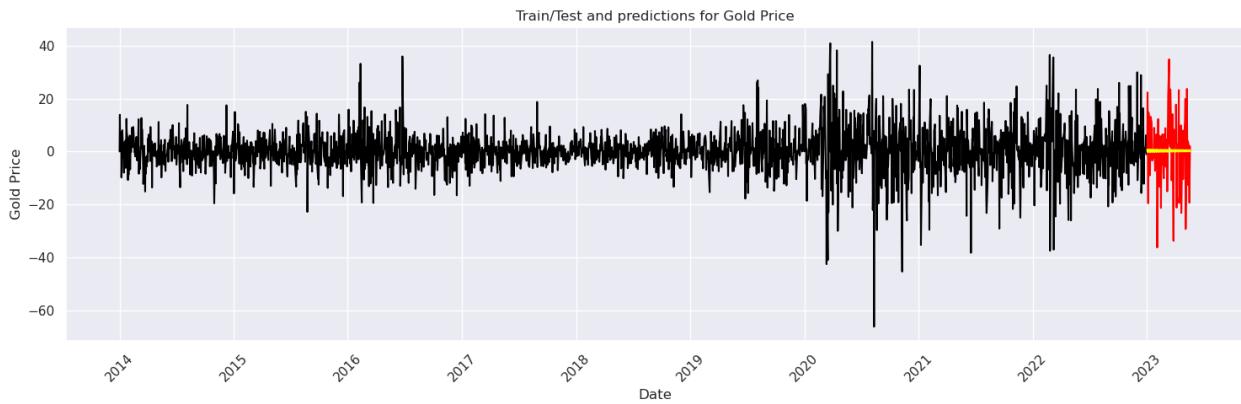
```
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)

SARIMAX Results
=====
Dep. Variable:      INR_diff    No. Observations:          2348
Model:             ARIMA(1, 1, 2)    Log Likelihood:       -8211.833
Date:        Mon, 29 May 2023   AIC:                   16431.665
Time:            14:03:08     BIC:                   16454.709
Sample:        01-01-2014   HQIC:                  16440.058
                    - 12-30-2022
Covariance Type: opg
=====
              coef    std err        z   P>|z|      [0.025]     [0.975]
-----+
ar.L1     -0.9939    0.011   -88.261   0.000    -1.016     -0.972
ma.L1     -0.0092    0.029    -0.319   0.750    -0.066     0.048
ma.L2     -0.9907    0.027   -36.776   0.000    -1.043     -0.938
sigma2    63.8507   1.986    32.156   0.000    59.959    67.743
=====
Ljung-Box (L1) (Q):      1.55   Jarque-Bera (JB):    3729.15
Prob(Q):                0.21   Prob(JB):                 0.00
Heteroskedasticity (H):  3.56   Skew:                  -0.22
Prob(H) (two-sided):    0.00   Kurtosis:                9.16
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

RMSE: 12.141379595326978





- After checking standard orders, we found that none of the orders are predicting the gold price well.

## Let's Try SARIMAX

### Helper function for checking order for SARIMAX

```
In [31]: def check_seasonal_order(p,d,q):
    y = dataIndia['INR_diff']
    from statsmodels.tsa.statespace.sarimax import SARIMAX
    SARIMAModel = SARIMAX(y, order = (p,d,q), seasonal_order = (0,0,0,12))
    ResultModel = SARIMAModel.fit()
    print(ResultModel.summary(), "\n")

    y_pred = ResultModel.get_forecast(len(test.index))
    y_pred_df = y_pred.conf_int(alpha = 0.05)
    y_pred_df[\"Predictions\"] = ResultModel.predict(start = y_pred_df.index[0], end = y_pred_df.index[-1])
    y_pred_df.index = test.index
    y_pred_out = y_pred_df[\"Predictions\"]

    from sklearn.metrics import mean_squared_error
    sarima_rmse = np.sqrt(mean_squared_error(test['INR_diff'].values, y_pred_df[\"Predictions\"]))
    print("RMSE: ", sarima_rmse, '\n')

    ResultModel.plot_diagnostics(figsize = (20,12))

    sns.set()
    plt.subplots(figsize = (15,10))
    plt.plot(train,color ='black',label = 'Train')
    plt.plot(test,color = 'red',label = 'Test')
    plt.plot(y_pred_out, color='green', label = 'ARIMA Predictions')
    plt.ylabel("Gold Price")
    plt.xlabel("Date")
    plt.xticks(rotation = 45)
    plt.title("Train/Test and predictions for Gold Price")
    plt.show()
```

```
In [32]: check_seasonal_order(2,1,2)

/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)
/opt/conda/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency B will be used.
  self._init_dates(dates, freq)
This problem is unconstrained.
```

## RUNNING THE L-BFGS-B CODE

\* \* \*

```

Machine precision = 2.220D-16
N =           5      M =        10

At X0          0 variables are exactly at the bounds

At iterate    0    f=  3.57609D+00  |proj g|=  5.90767D-02
At iterate    5    f=  3.53466D+00  |proj g|=  2.25843D-02
At iterate   10    f=  3.52324D+00  |proj g|=  8.20462D-04
At iterate   15    f=  3.52321D+00  |proj g|=  6.80598D-04
At iterate   20    f=  3.52317D+00  |proj g|=  1.76131D-04

```

\* \* \*

```

Tit  = total number of iterations
Tnf  = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

\* \* \*

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
5	24	29	1	0	0	1.402D-05	3.523D+00
F = 3.5231662851013223							

CONVERGENCE: REL\_REDUCTION\_OF\_F\_<=\_FACTR\*EPSMCH  
SARIMAX Results

```
=====
Dep. Variable:           INR_diff    No. Observations:             2448
Model:                 SARIMAX(2, 1, 2)    Log Likelihood       -8624.711
Date:                 Mon, 29 May 2023   AIC                  17259.422
Time:                 14:03:13         BIC                  17288.435
Sample:                01-01-2014   HQIC                  17269.967
                           - 05-19-2023
Covariance Type:            opg
=====
```

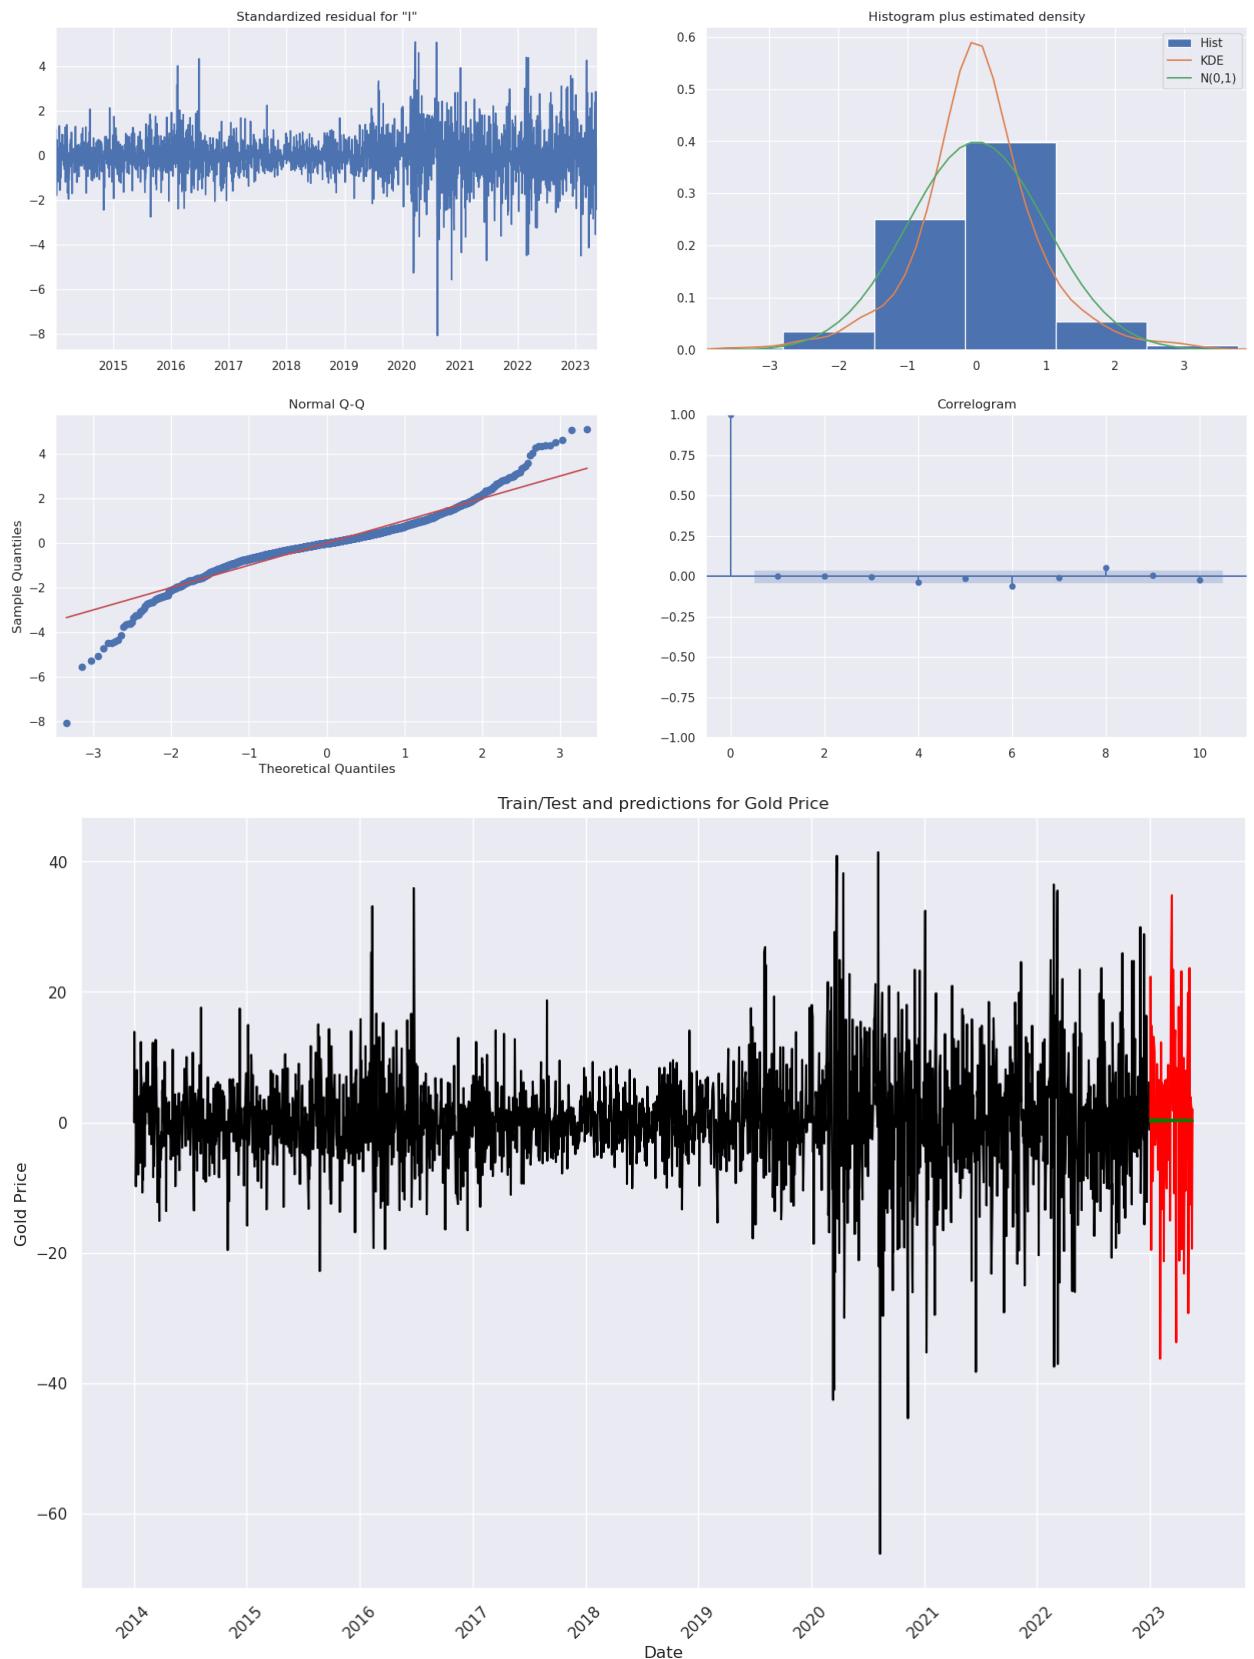
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-1.0202	0.024	-42.217	0.000	-1.068	-0.973
ar.L2	-0.0261	0.015	-1.700	0.089	-0.056	0.004
ma.L1	-0.0077	0.024	-0.322	0.748	-0.055	0.039
ma.L2	-0.9922	0.022	-45.446	0.000	-1.035	-0.949
sigma2	67.2326	1.555	43.239	0.000	64.185	70.280

```
=====
Ljung-Box (L1) (Q):            0.00    Jarque-Bera (JB):        3369.79
Prob(Q):                      0.99    Prob(JB):                   0.00
Heteroskedasticity (H):        3.42    Skew:                     -0.23
Prob(H) (two-sided):           0.00    Kurtosis:                  8.73
=====
```

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

RMSE: 12.135846463698607



- SARIMAX is also not predicting well as it was expected

## Let's Try Prophet

### Model Selection

```
In [33]: dataIndia = dataIndia.reset_index().drop(columns = 'INR_diff')
dataIndia.head()
```

```
Out[33]:
```

	Date	Year	Month	Week	Day	Indian rupee
0	2014-01-01	2014		1	1	609.17
1	2014-01-02	2014		1	2	623.05
2	2014-01-03	2014		1	1	627.90
3	2014-01-06	2014		1	2	635.28
4	2014-01-07	2014		1	2	625.47

```
In [34]:
```

```
from sklearn.model_selection import train_test_split
x = dataIndia[['Year','Month','Week','Day']]
y = dataIndia[['Indian rupee']]
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.2, random_state = 99)
```

```
In [35]:
```

```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_train,y_train)
y_pred = model.predict(x_test)
from sklearn.metrics import r2_score
r2_score(y_test,y_pred)
```

```
Out[35]: 0.8415031930138567
```

```
In [36]: model.score(x_train,y_train)
```

```
Out[36]: 0.8493969510646942
```

- There are no overfitting issues, but let's find the best model with the best accuracy

## Importing models

```
In [37]:
```

```
from sklearn.tree import DecisionTreeRegressor
dtr = DecisionTreeRegressor()
from sklearn.ensemble import RandomForestRegressor
rfr = RandomForestRegressor()
from xgboost import XGBRegressor
xgb = XGBRegressor()
from lightgbm import LGBMRegressor
lgbm = LGBMRegressor()
```

```
In [38]:
```

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
```

## Helper function for model evaluation

```
In [39]:
```

```
def evaluation_metrics(model):
    model.fit(x_train,y_train)
    y_pred = model.predict(x_test)
    mse = mean_squared_error(y_test,y_pred)
    rmse = np.sqrt(mse)
    mae = mean_absolute_error(y_test,y_pred)
    r2score = r2_score(y_test,y_pred)
    trainscore = model.score(x_train,y_train)

    print(f'Mean Squared Error = {mse}')
    print(f'Mean Absolute Error = {mae}')
    print(f'Root Mean Squared Error = {rmse}')
    print(f'r2 score = {r2score}')

    Comparison = x_test.copy("deep")
    Comparison['Actual Price'] = y_test
    Comparison['Predicted Price'] = y_pred

    sns.lineplot(Comparison, x = 'Year',y = 'Actual Price')
    sns.lineplot(Comparison, x = 'Year',y = 'Predicted Price')
    plt.legend(['Actual Price','Predicted Price'])
    plt.show()

    return {'Training Score': trainscore,
            'Testing Score': r2score}
```

```
In [40]: results = {}
```

## Linear Regression

```
In [41]:
```

```
results['Linear Regression'] = evaluation_metrics(model)

Mean Squared Error = 7985.63573389267
Mean Absolute Error = 74.62807807771928
Root Mean Squared Error = 89.36238433419662
r2 score = 0.8415031930138567
```



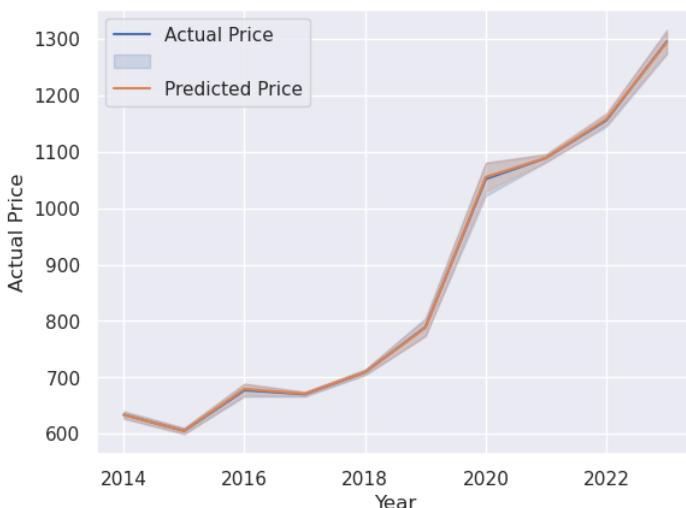
## Decision Tree Regressor

```
In [42]: results['Decision Tree Regressor'] = evaluation_metrics(dtr)
Mean Squared Error = 76.2006957142857
Mean Absolute Error = 5.906591836734695
Root Mean Squared Error = 8.729300986578805
r2 score = 0.9984875885448197
```



## Random Forest Regressor

```
In [43]: results['Random Forest Regressor'] = evaluation_metrics(rfr)
Mean Squared Error = 70.18512658581592
Mean Absolute Error = 5.604885918367294
Root Mean Squared Error = 8.377656389815467
r2 score = 0.9986069839856886
```



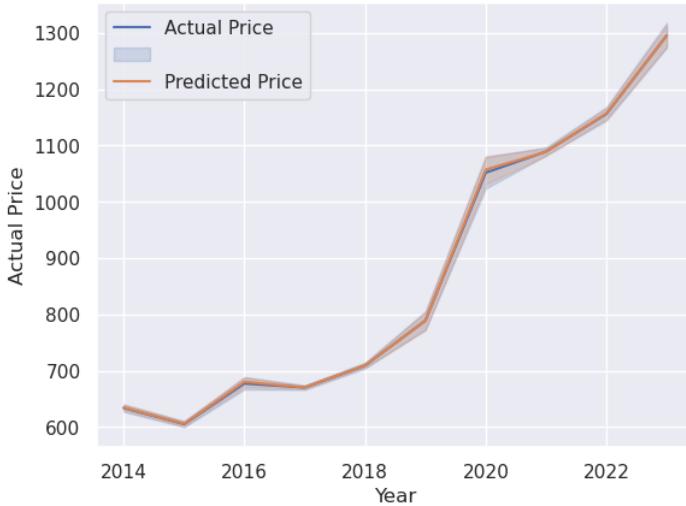
## XGB Regressor

```
In [44]: results['XGB Regressor'] = evaluation_metrics(xgb)
Mean Squared Error = 115.93633434047514
Mean Absolute Error = 7.125565788424744
Root Mean Squared Error = 10.767373604573919
r2 score = 0.9976989259942507
```



## LGBM Regressor

```
In [45]: results['LGBM Regressor'] = evaluation_metrics(lgbm)
Mean Squared Error = 132.8857599566127
Mean Absolute Error = 8.35661415136721
Root Mean Squared Error = 11.527608511554392
r2 score = 0.9973625182717947
```



## Model comparison

```
In [46]: resultsdf = pd.DataFrame(results)
resultsdf
```

```
Out[46]:
```

	Linear Regression	Decision Tree Regressor	Random Forest Regressor	XGB Regressor	LGBM Regressor
Training Score	0.849397	1.000000	0.999789	0.999713	0.998575
Testing Score	0.841503	0.998488	0.998607	0.997699	0.997363

- Random Forest Regressor is best in all of the models

```
In [47]: pip install prophet
```

```

Requirement already satisfied: prophet in /opt/conda/lib/python3.7/site-packages (1.1.1)
Requirement already satisfied: setuptools>=42 in /opt/conda/lib/python3.7/site-packages (from prophet) (59.8.0)
Requirement already satisfied: holidays>=0.14.2 in /opt/conda/lib/python3.7/site-packages (from prophet) (0.21.13)
Requirement already satisfied: tqdm>=4.36.1 in /opt/conda/lib/python3.7/site-packages (from prophet) (4.64.1)
Requirement already satisfied: pandas>=1.0.4 in /opt/conda/lib/python3.7/site-packages (from prophet) (1.3.5)
Requirement already satisfied: matplotlib>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from prophet) (3.5.3)
Requirement already satisfied: cmdstanpy>=1.0.4 in /opt/conda/lib/python3.7/site-packages (from prophet) (1.1.0)
Requirement already satisfied: convertdate>=2.1.2 in /opt/conda/lib/python3.7/site-packages (from prophet) (2.4.0)
Requirement already satisfied: setuptools-git>=1.2 in /opt/conda/lib/python3.7/site-packages (from prophet) (1.2)
Requirement already satisfied: python-dateutil>=2.8.0 in /opt/conda/lib/python3.7/site-packages (from prophet) (2.8.2)
Requirement already satisfied: wheel>=0.37.0 in /opt/conda/lib/python3.7/site-packages (from prophet) (0.38.4)
Requirement already satisfied: LunarCalendar>=0.0.9 in /opt/conda/lib/python3.7/site-packages (from prophet) (0.0.9)
Requirement already satisfied: numpy>=1.15.4 in /opt/conda/lib/python3.7/site-packages (from prophet) (1.21.6)
Requirement already satisfied: pomegranate<1,>0.3.13 in /opt/conda/lib/python3.7/site-packages (from convertdate>=2.1.2->prophet) (0.5.12)
Requirement already satisfied: backports.zoneinfo in /opt/conda/lib/python3.7/site-packages (from holidays>=0.14.2->prophet) (0.2.1)
Requirement already satisfied: korean-lunar-calendar in /opt/conda/lib/python3.7/site-packages (from holidays>=0.14.2->prophet) (0.3.1)
Requirement already satisfied: hijri-converter in /opt/conda/lib/python3.7/site-packages (from holidays>=0.14.2->prophet) (2.2.4)
Requirement already satisfied: ephem>=3.7.5.3 in /opt/conda/lib/python3.7/site-packages (from LunarCalendar>=0.0.9->prophet) (4.1.4)
Requirement already satisfied: pytz in /opt/conda/lib/python3.7/site-packages (from LunarCalendar>=0.0.9->prophet) (2022.7.1)
Requirement already satisfied: fonttools>=4.22.0 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=2.0.0->prophet) (4.38.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=2.0.0->prophet) (1.4.4)
Requirement already satisfied: pyparsing>=2.2.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=2.0.0->prophet) (3.0.9)
Requirement already satisfied: pillow>=6.2.0 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=2.0.0->prophet) (9.4.0)
Requirement already satisfied: cycler>=0.10 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=2.0.0->prophet) (0.11.0)
Requirement already satisfied: packaging>=20.0 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=2.0.0->prophet) (23.0)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.7/site-packages (from python-dateutil>=2.8.0->prophet) (1.16.0)
Requirement already satisfied: typing-extensions in /opt/conda/lib/python3.7/site-packages (from kiwisolver>=1.0.1>matplotlib>=2.0.0->prophet) (4.4.0)
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv
Note: you may need to restart the kernel to use updated packages.

```

## Making time dataframe for Prophet

```
In [48]: time_data = dataIndia[['Date','Indian rupee']]
time_data.columns = ['ds','y']
time_data.head()
```

```
Out[48]:      ds      y
0 2014-01-01  609.17
1 2014-01-02  623.05
2 2014-01-03  627.90
3 2014-01-06  635.28
4 2014-01-07  625.47
```

```
In [49]: from prophet import Prophet
m = Prophet(daily_seasonality=True)
model_fit = m.fit(time_data)
```

```
14:03:46 - cmdstanpy - INFO - Chain [1] start processing
14:03:48 - cmdstanpy - INFO - Chain [1] done processing
```

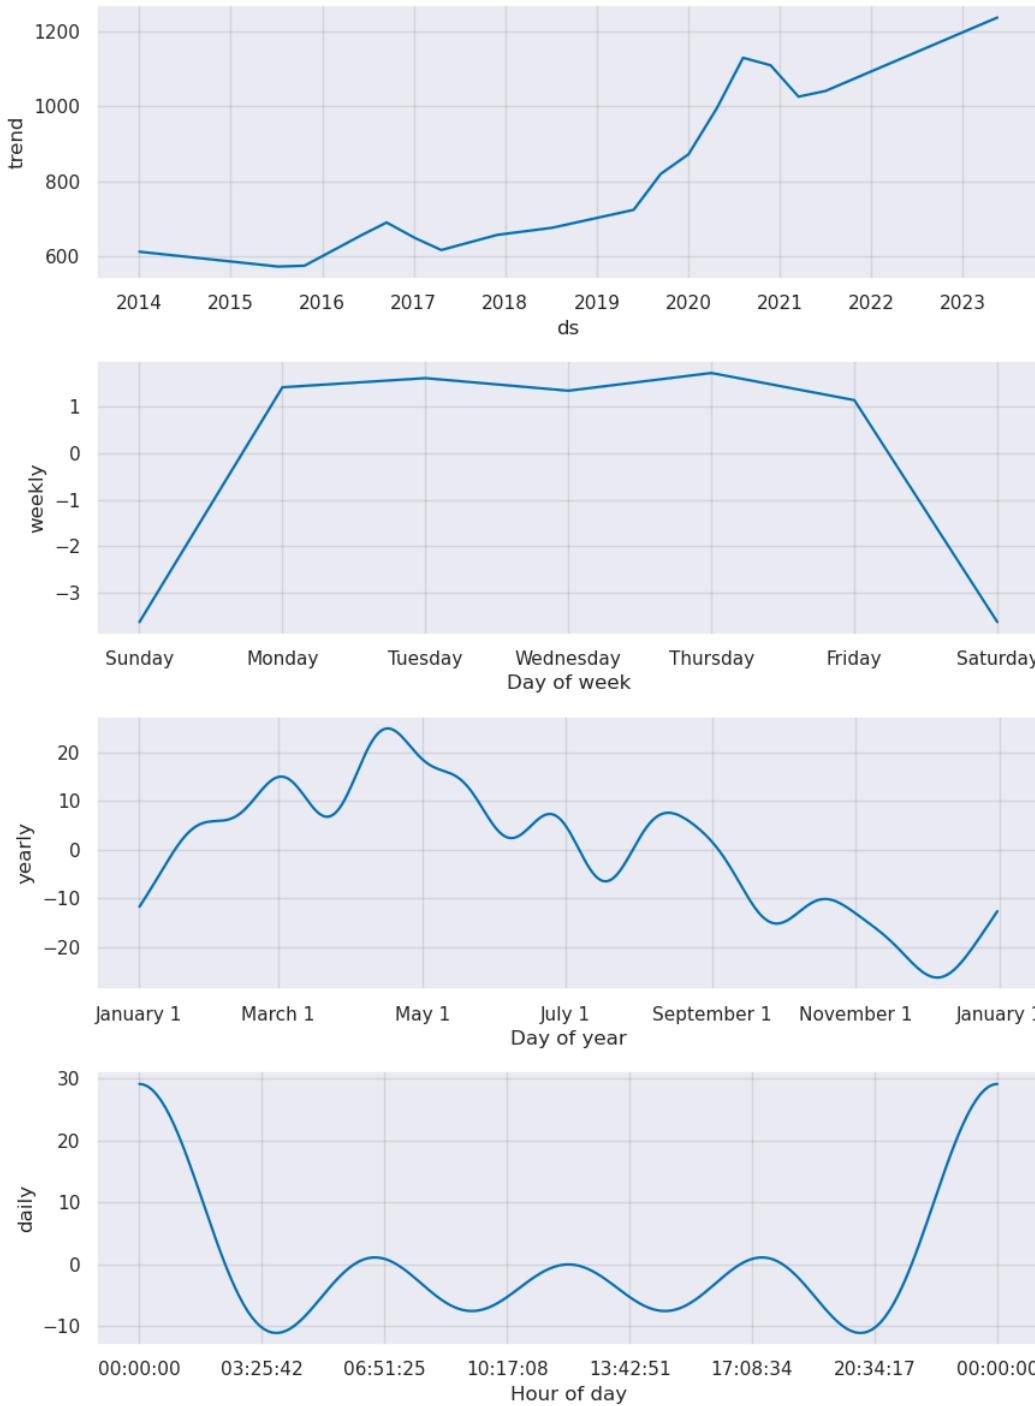
```
In [50]: forecast = model_fit.predict(time_data)
forecast[['ds','yhat','yhat_upper','yhat_lower']].tail()
```

```
Out[50]:      ds      yhat    yhat_upper    yhat_lower
2443 2023-05-15  1281.642774  1315.635540  1247.050731
2444 2023-05-16  1281.841554  1316.169976  1247.779814
2445 2023-05-17  1281.524603  1316.686618  1244.436156
2446 2023-05-18  1281.806680  1315.260755  1247.167047
2447 2023-05-19  1281.064695  1318.194072  1247.558369
```

```
In [51]: model_fit.plot(forecast);
```



```
In [52]: model_fit.plot_components(forecast);
```



Month of June and August are awesome

Forecasting for next 6 months

```
In [53]: m2 = Prophet(daily_seasonality=True)
m2.fit(time_data)
future = m2.make_future_dataframe(periods = 336)
future.tail(10)

14:03:54 - cmdstanpy - INFO - Chain [1] start processing
14:03:56 - cmdstanpy - INFO - Chain [1] done processing
```

```
Out[53]:      ds
2774 2024-04-10
2775 2024-04-11
2776 2024-04-12
2777 2024-04-13
2778 2024-04-14
2779 2024-04-15
2780 2024-04-16
2781 2024-04-17
2782 2024-04-18
2783 2024-04-19
```

```
In [54]: forecast2 = m2.predict(future)
forecast2[['ds','yhat','yhat_upper','yhat_lower']].tail()
```

```
Out[54]:      ds    yhat  yhat_upper  yhat_lower
2779 2024-04-15  1387.098424  1513.910925  1247.925641
2780 2024-04-16  1387.651489  1507.342940  1250.645201
2781 2024-04-17  1387.623199  1513.693186  1243.390041
2782 2024-04-18  1388.140714  1517.181194  1248.928742
2783 2024-04-19  1387.593558  1514.821713  1250.925275
```

```
In [55]: m2.plot(forecast2);
```



- In the starting of 2024, the gold price will inflate

```
In [56]: m2.plot_components(forecast2);
```



```
In [57]: predictions = forecast2[['ds','yhat']]
predictions = predictions.rename(columns = {'ds':'Date','yhat':'Indian rupee'})
predictions.head()
```

```
Out[57]:
```

	Date	Indian rupee
0	2014-01-01	630.941574
1	2014-01-02	631.996564
2	2014-01-03	632.092243
3	2014-01-06	634.467958
4	2014-01-07	635.382572

```
In [58]: plt.subplots(figsize = (15,8))
sns.lineplot(time_data,x = 'ds',y='y',legend='full')
sns.lineplot(predictions,x = 'Date',y = 'Indian rupee',legend='full')
plt.legend(labels = ['Actual','','Prediction'])
plt.show()
```



```
In [59]: predictions['Date'] = pd.to_datetime(predictions['Date'])
predictions['Year'] = predictions['Date'].dt.year
predictions['Month'] = predictions['Date'].dt.month
predictions['Week'] = predictions['Date'].dt.week
predictions['Day'] = predictions['Date'].dt.day
```

## Model Evaluation

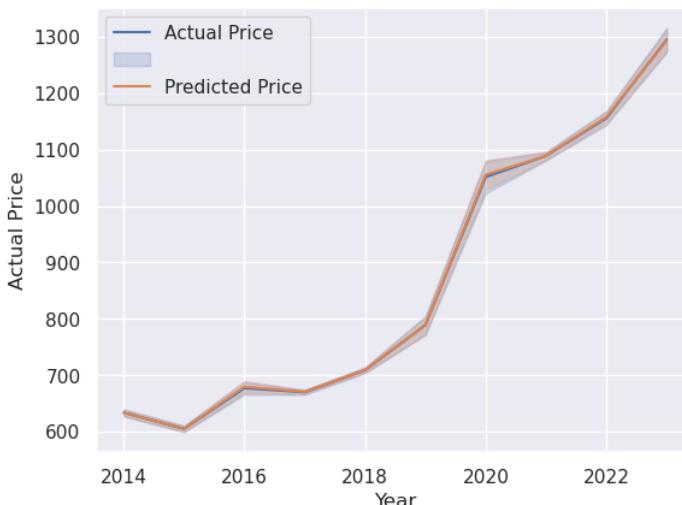
```
In [60]: x_time = predictions.drop(columns=['Date', 'Indian rupee'])
y_time = predictions[['Indian rupee']]
```

```
In [61]: x_time_train,x_time_test,y_time_train,y_time_test = train_test_split(x_time,y_time,test_size=0.2,random_state=99)
```

```
In [62]: rfr_time_model = RandomForestRegressor()
rfr_time_model.fit(x_time_train,y_time_train)
rfr_time_pred = rfr_time_model.predict(x_time_test)
```

```
In [63]: results['Random Forest Evaluation Regressor'] = evaluation_metrics(rfr_time_model)

Mean Squared Error = 68.51278588663261
Mean Absolute Error = 5.57226306122445
Root Mean Squared Error = 8.277245066242306
r2 score = 0.9986401761659791
```



```
In [64]: del results['Linear Regression']
del results['Decision Tree Regressor']
del results['XGB Regressor']
del results['LGBM Regressor']
```

## Comparing accuracy of predicted model and forecasted model

```
In [65]: CompareRFR = pd.DataFrame(results)
```

Out[65]:

	Random Forest Regressor	Random Forest Evaluation Regressor
Training Score	0.999789	0.999805
Testing Score	0.998607	0.998640

What If !!! We take the whole data (1979 - 2023) : Will it be better forecasting?

```
In [66]: dataIndiaAll = data[['Date','Year','Month','Week','Day','Indian rupee']]
dataIndiaAll = dataIndiaAll.dropna().reset_index(drop=True)
dataIndiaAll.head()
```

Out[66]:

	Date	Year	Month	Week	Day	Indian rupee
0	1979-01-02	1979	1	1	2	14.66
1	1979-01-03	1979	1	1	3	14.73
2	1979-01-04	1979	1	1	4	14.81
3	1979-01-05	1979	1	1	5	15.07
4	1979-01-08	1979	1	2	8	15.06

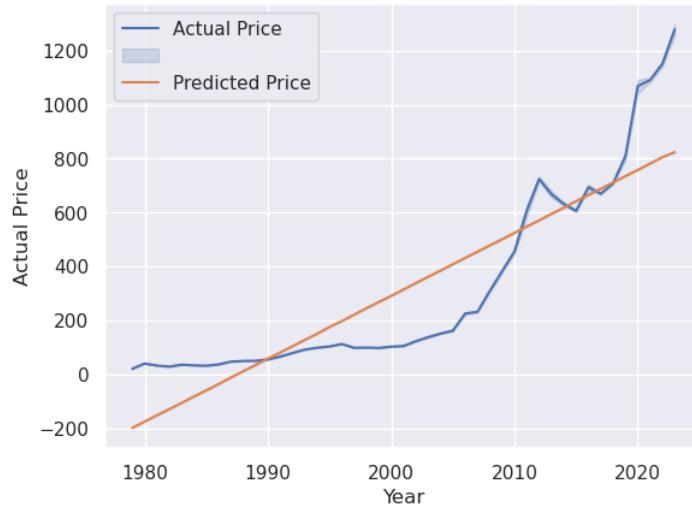
## Model selection for whole data

```
In [67]: x_all = dataIndiaAll[['Year','Month','Week','Day']]
y_all = dataIndiaAll[['Indian rupee']]
x_train,x_test,y_train,y_test = train_test_split(x_all,y_all,test_size = 0.2, random_state = 99)
```

```
In [68]: AllDataResults = {}
```

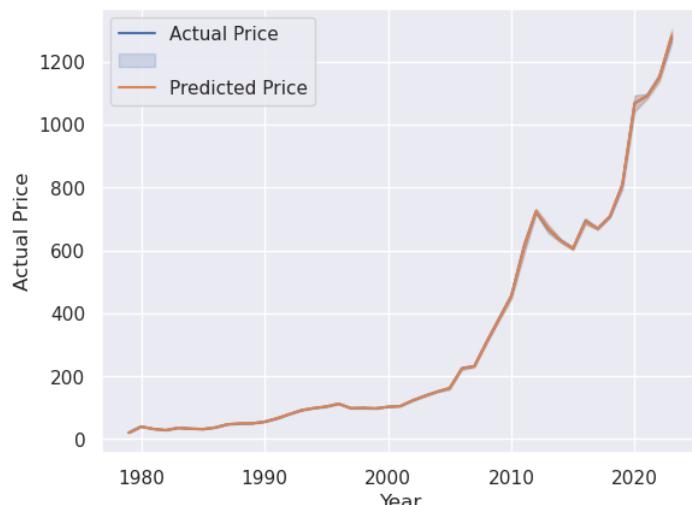
```
In [69]: AllDataResults['Linear Regression'] = evaluation_metrics(model)
```

```
Mean Squared Error = 26197.346839378588
Mean Absolute Error = 129.9110398940336
Root Mean Squared Error = 161.85594471436193
r2 score = 0.7648229167312672
```



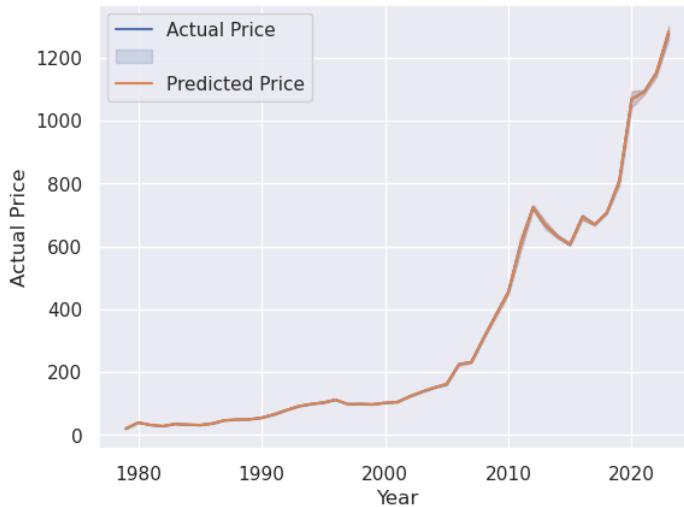
```
In [70]: AllDataResults['Decision Tree Regressor'] = evaluation_metrics(dtr)
```

```
Mean Squared Error = 20.869293911917087
Mean Absolute Error = 2.3210060449050083
Root Mean Squared Error = 4.568292231448979
r2 score = 0.9998126535598366
```



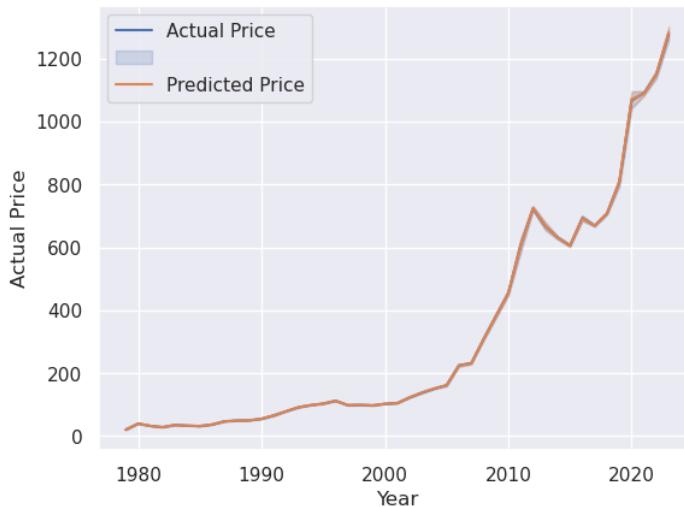
```
In [71]: AllDataResults['Random Forest Regressor'] = evaluation_metrics(rfr)
```

```
Mean Squared Error = 17.457940534645775  
Mean Absolute Error = 2.156791796200346  
Root Mean Squared Error = 4.178270040895606  
r2 score = 0.9998432777349557
```



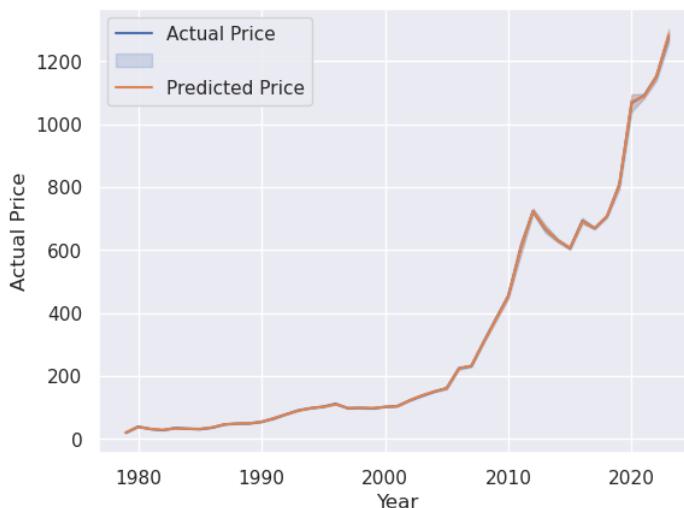
```
In [72]: AllDataResults['XGB Regressor'] = evaluation_metrics(xgb)
```

```
Mean Squared Error = 27.853747351798226  
Mean Absolute Error = 3.096116317656785  
Root Mean Squared Error = 5.2776649525901345  
r2 score = 0.9997499531879902
```



```
In [73]: AllDataResults['LGBM Regressor'] = evaluation_metrics(lgbm)
```

```
Mean Squared Error = 53.66628227226071  
Mean Absolute Error = 4.6011194030607  
Root Mean Squared Error = 7.325727422738352  
r2 score = 0.9995182306127391
```



```
In [74]: AllDataResultsdf = pd.DataFrame(AllDataResults)
```

AllDataResultsdf																		
Out[74]:																		
<table> <thead> <tr> <th></th><th>Linear Regression</th><th>Decision Tree Regressor</th><th>Random Forest Regressor</th><th>XGB Regressor</th><th>LGBM Regressor</th></tr> </thead> <tbody> <tr> <td>Training Score</td><td>0.767883</td><td>1.000000</td><td>0.999969</td><td>0.999885</td><td>0.999539</td></tr> <tr> <td>Testing Score</td><td>0.764823</td><td>0.999813</td><td>0.999843</td><td>0.999750</td><td>0.999518</td></tr> </tbody> </table>		Linear Regression	Decision Tree Regressor	Random Forest Regressor	XGB Regressor	LGBM Regressor	Training Score	0.767883	1.000000	0.999969	0.999885	0.999539	Testing Score	0.764823	0.999813	0.999843	0.999750	0.999518
	Linear Regression	Decision Tree Regressor	Random Forest Regressor	XGB Regressor	LGBM Regressor													
Training Score	0.767883	1.000000	0.999969	0.999885	0.999539													
Testing Score	0.764823	0.999813	0.999843	0.999750	0.999518													
• Random Forest Regressor is performing best, but LGBM Regressor is performing very accurate to 4 decimals.																		

## Forecasting whole data

```
In [75]: time_data_all = dataIndiaAll[['Date','Indian rupee']]
time_data_all.columns = ['ds','y']
time_data_all.head()
```

```
Out[75]:
```

	ds	y
0	1979-01-02	14.66
1	1979-01-03	14.73
2	1979-01-04	14.81
3	1979-01-05	15.07
4	1979-01-08	15.06

```
In [76]: m3 = Prophet(daily_seasonality=True)
m3.fit(time_data_all)

14:04:32 - cmdstanpy - INFO - Chain [1] start processing
14:04:45 - cmdstanpy - INFO - Chain [1] done processing
```

```
Out[76]: <prophet.forecaster.Prophet at 0x7e051132c650>
```

```
In [77]: forecast3 = m3.predict(time_data_all)
forecast3[['ds','yhat','yhat_upper','yhat_lower']].tail()
```

```
Out[77]:
```

	ds	yhat	yhat_upper	yhat_lower
11574	2023-05-15	1211.003923	1261.959662	1156.653545
11575	2023-05-16	1211.259981	1268.571187	1158.375886
11576	2023-05-17	1211.281962	1261.144123	1158.656924
11577	2023-05-18	1211.419061	1261.241815	1156.857749
11578	2023-05-19	1211.357644	1263.998107	1157.240909



```
In [79]: m3.plot_components(forecast3);
```



This is different: April, June, and July are best month here

```
In [80]: m4 = Prophet(daily_seasonality=True)
m4.fit(time_data_all)
future2 = m4.make_future_dataframe(periods = 336)
future2.tail(10)
```

```
14:04:54 - cmdstanpy - INFO - Chain [1] start processing
14:05:08 - cmdstanpy - INFO - Chain [1] done processing
```

```
Out[80]: ds
11905 2024-04-10
11906 2024-04-11
11907 2024-04-12
11908 2024-04-13
11909 2024-04-14
11910 2024-04-15
11911 2024-04-16
11912 2024-04-17
11913 2024-04-18
11914 2024-04-19
```

```
In [81]: forecast4 = m4.predict(future2)
forecast4[['ds','yhat','yhat_upper','yhat_lower']].tail()
```

```
Out[81]:      ds    yhat  yhat_upper  yhat_lower
11910 2024-04-15 1284.354958 1341.617513 1232.523239
11911 2024-04-16 1284.696916 1343.013181 1227.940609
11912 2024-04-17 1284.807629 1344.774432 1229.777855
11913 2024-04-18 1285.037346 1343.159418 1227.441624
11914 2024-04-19 1285.073246 1343.482370 1231.961452
```

```
In [82]: m4.plot(forecast4);
```



```
In [83]: m4.plot_components(forecast4);
```



```
In [84]: predictions2 = forecast4[['ds','yhat']]
predictions2 = predictions2.rename(columns = {'ds':'Date','yhat':'Indian rupee'})
predictions2.head()
```

	Date	Indian rupee
0	1979-01-02	18.571801
1	1979-01-03	18.717701
2	1979-01-04	18.997408
3	1979-01-05	19.094261
4	1979-01-08	19.919275

```
In [85]: plt.subplots(figsize = (15,8))
sns.lineplot(time_data_all,x = 'ds',y='y',legend='full')
sns.lineplot(predictions2,x = 'Date',y = 'Indian rupee',legend='full')
plt.legend(labels = ['Actual','','Prediction'])
plt.show()
```



Doesn't looks good

```
In [86]: predictions2['Date'] = pd.to_datetime(predictions2['Date'])
predictions2['Year'] = predictions2['Date'].dt.year
predictions2['Month'] = predictions2['Date'].dt.month
predictions2['Week'] = predictions2['Date'].dt.week
predictions2['Day'] = predictions2['Date'].dt.day
```

## Model Evaluation for whole data

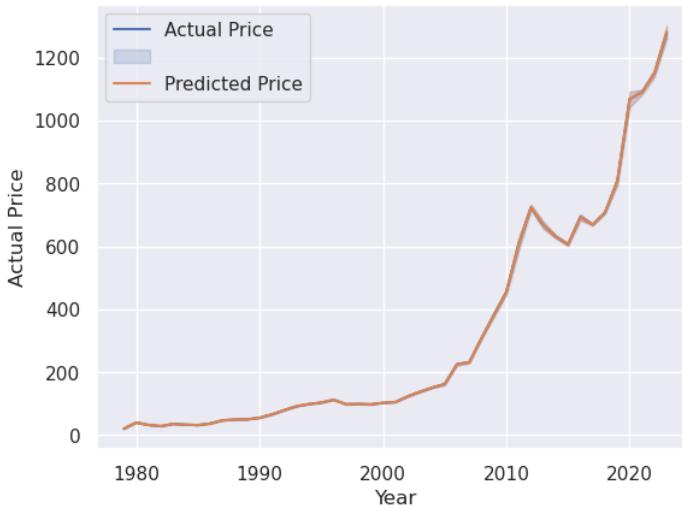
```
In [87]: x_time_all = predictions2.drop(columns=['Date', 'Indian rupee'])
y_time_all = predictions2[['Indian rupee']]
```

```
In [88]: x_time_train,x_time_test,y_time_train,y_time_test = train_test_split(x_time_all,y_time_all,test_size=0.2,random_state=99)
```

```
In [89]: rfr_time_all_model = RandomForestRegressor()
rfr_time_all_model.fit(x_time_train,y_time_train)
rfr_time_all_pred = rfr_time_all_model.predict(x_time_test)
```

```
In [90]: AllDataResults['Random Forest Evaluation Regressor'] = evaluation_metrics(rfr_time_all_model)
```

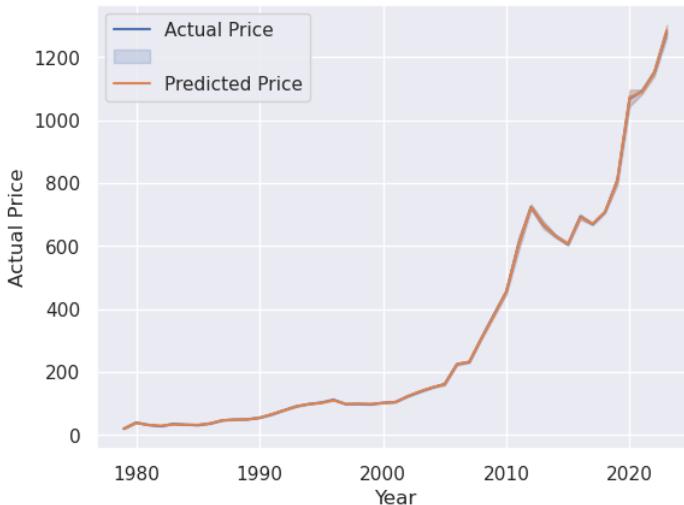
```
Mean Squared Error = 17.095929741446355
Mean Absolute Error = 2.142367227979272
Root Mean Squared Error = 4.134722450352182
r2 score = 0.999846527554221
```



```
In [91]: lgbm_time_all_model = LGBMRegressor()
lgbm_time_all_model.fit(x_time_train,y_time_train)
lgbm_time_all_pred = lgbm_time_all_model.predict(x_time_test)
```

```
In [92]: AllDataResults['LGBM Evaluation Regressor'] = evaluation_metrics(lgbm_time_all_model)
```

```
Mean Squared Error = 53.66628227226071
Mean Absolute Error = 4.6011194030607
Root Mean Squared Error = 7.325727422738352
r2 score = 0.9995182306127391
```



```
In [93]: del AllDataResults['Linear Regression']
del AllDataResults['Decision Tree Regressor']
del AllDataResults['XGB Regressor']
```

### Comparing selected and predicted models for the whole data

```
In [94]: CompareAll = pd.DataFrame(AllDataResults)
CompareAll
```

	Random Forest Regressor	LGBM Regressor	Random Forest Evaluation Regressor	LGBM Evaluation Regressor
Training Score	0.999969	0.999539	0.999970	0.999539
Testing Score	0.999843	0.999518	0.999847	0.999518

- LGBM Regressor model is very very accurate as it was before the forecasting, but Random Forest Regressor has the highest accuracy

### Looking the difference between 10 years prediction and the whole data prediction

```
In [95]: analysis1 = predictions[predictions['Year']>2022]
analysis2 = predictions2[predictions2['Year']>2022]

plt.subplots(figsize = (15,5))
plt.plot(analysis1['Date'],analysis1['Indian rupee'],color = 'red',label = 'Predictions from 10 years')
plt.plot(analysis2['Date'],analysis2['Indian rupee'],color = 'blue',label = 'Predictions from all available data')
plt.legend()
plt.show()
```



### Conclusion:

1. Always large data doesn't forecast good, however model accuracy is pretty good as in LGBM Regressor.
2. The common month of low gold price in both the forecasting is June.
3. Compared to the monthly forecast of the 10 years past data, first forecasting results are good. But, the second forecasting results does not match any of the trends of the whole data forecasting components plot. Clearly, the whole data did not forecasted well.
4. The real way of evaluating this forecasting to see the gold price in June and August. Exciting!!!